
Interrupts

Concurrency

Three Classic forms of Concurrency

Multiprocessing

Using two or more CPUs in order to execute processes simultaneously.

Time-Sharing

Sharing a given CPU resource amongst multiple processes via a protocol for sharing execution time. Can be *preemptive* or *cooperative*.

Interrupts

Hardware interrupts the CPU under some circumstances, providing an opportunity to switch processes.

Interrupts

- An **interrupt** is a signal asking for CPU attention
 - CPU forces a jump to the **interrupt handler**
 - Like a function call
 - Can happen from hardware or software
- Original process doesn't run again until handler returns (maybe)

Interrupts

- If you don't **disable interrupts**, any program can be interrupted (user or kernel)
- An interrupt must be able to run on top of **any** other process
- Thus, interrupts provide a mechanism to build pseudo-concurrency on top of

Timer Interrupts Visual Example

Proc A

```
mov    %rax, -0x8(%rbp)
xor     %eax, %eax
lea     -0x1c(%rbp), %rax
mov     %rax, %rdi
```

CPU

TIMER

Interrupt Vector

0x0057

0x0181

0x0237

0x0302

....

Proc A

```
mov    %rax, -0x8(%rbp)
xor     %eax, %eax
lea     -0x1c(%rbp), %rax
mov     %rax, %rdi
```

Interrupt Vector

0x0057

0x0181

0x0237

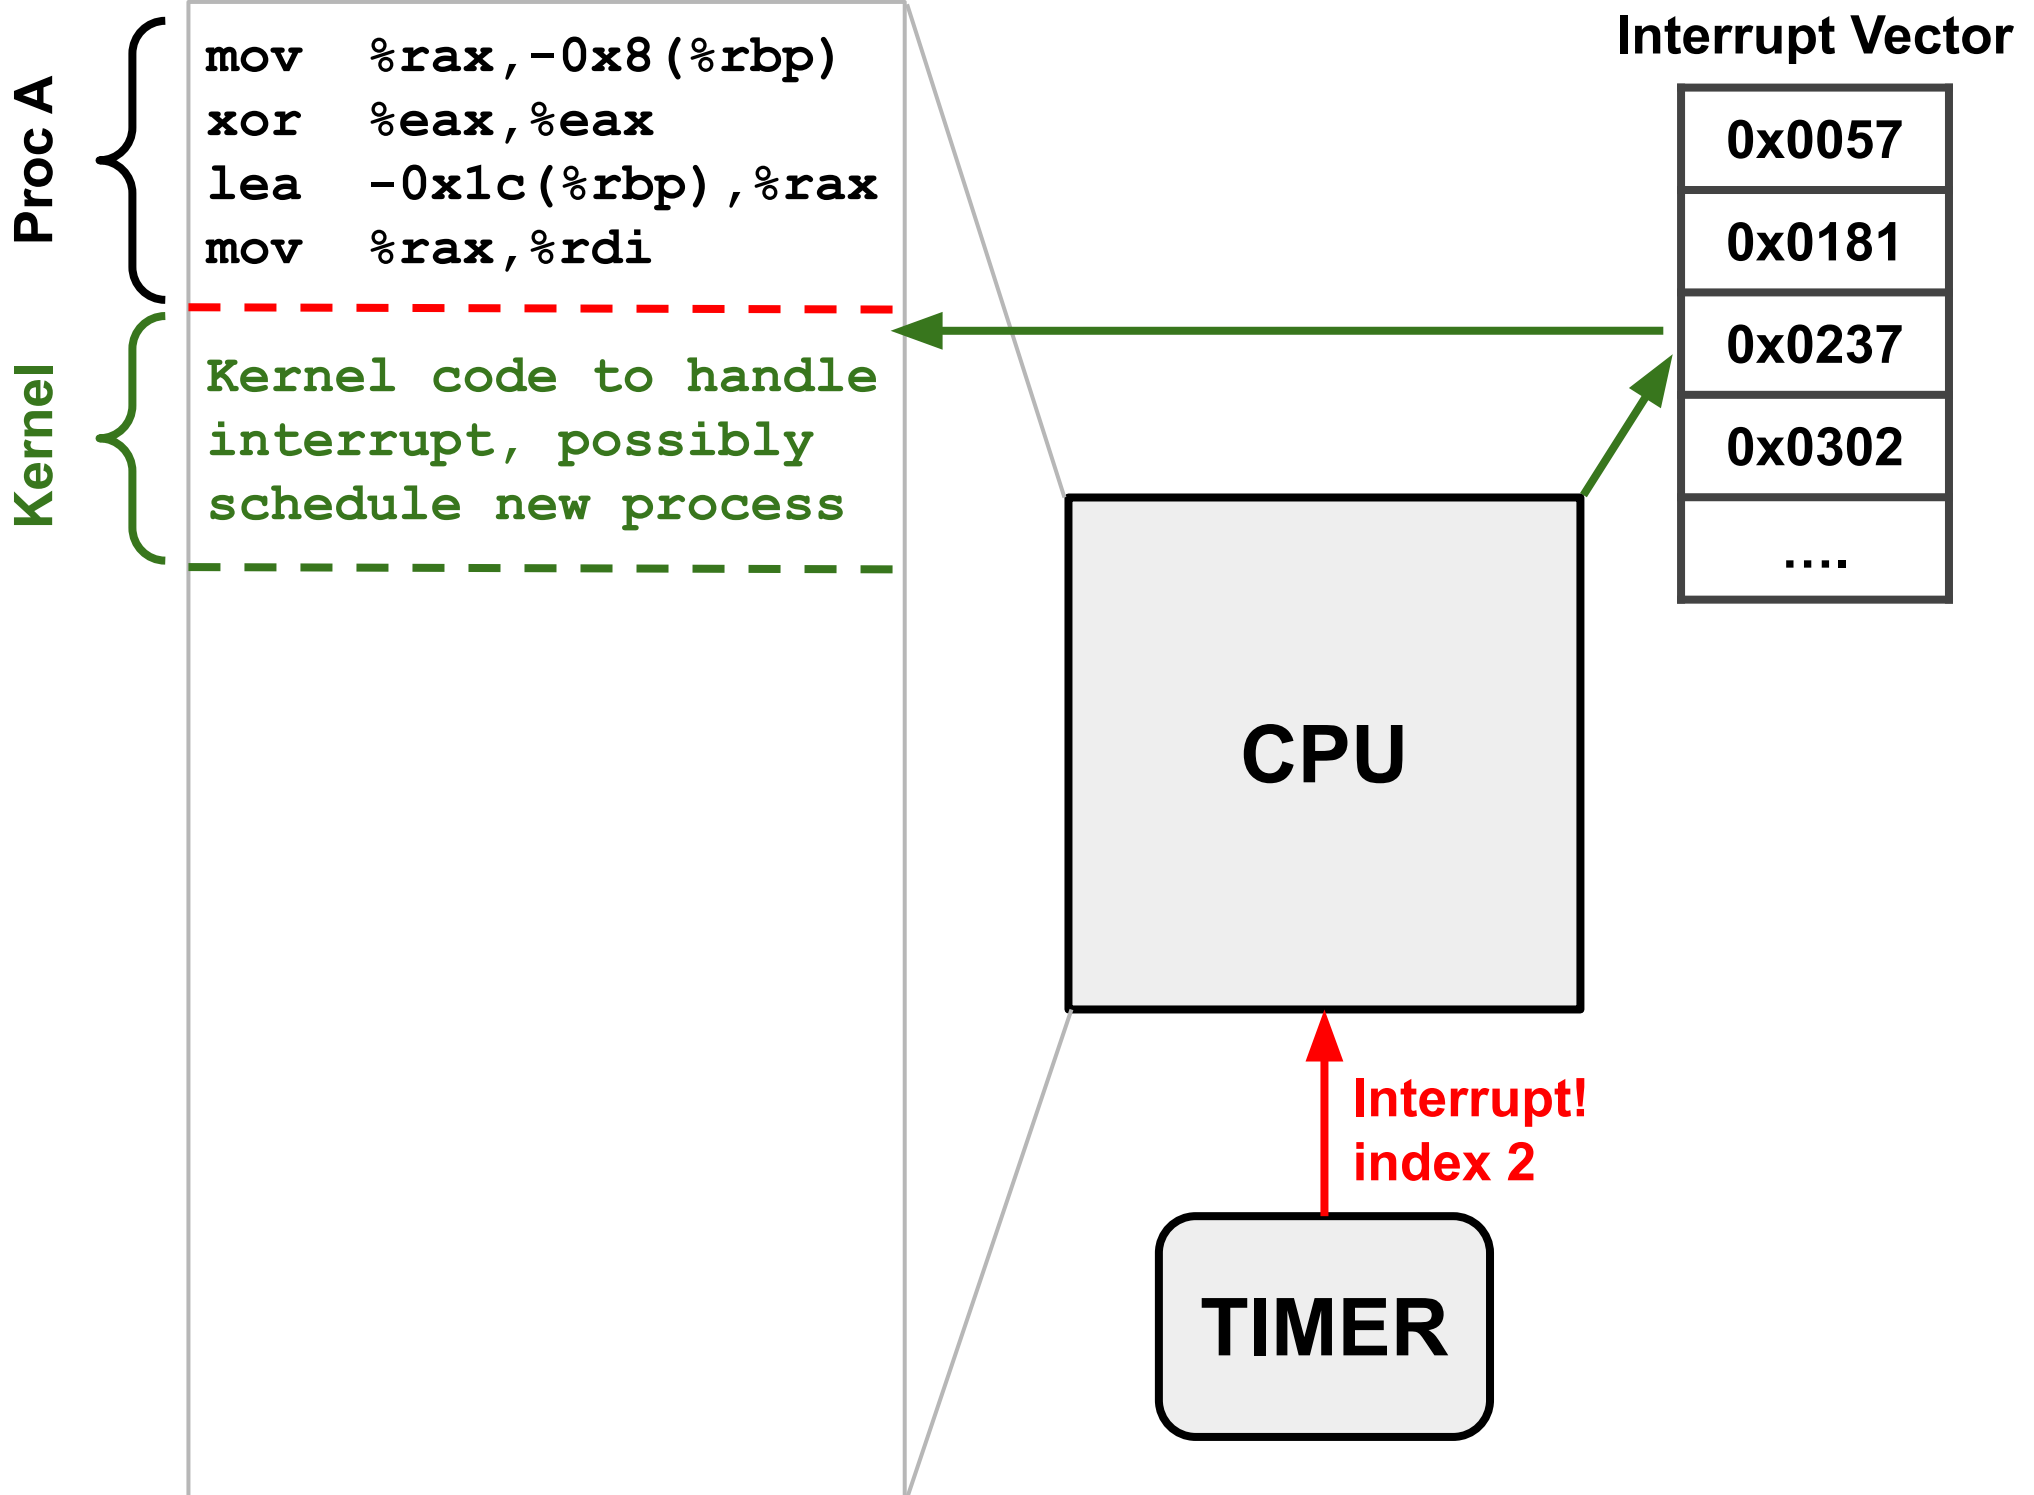
0x0302

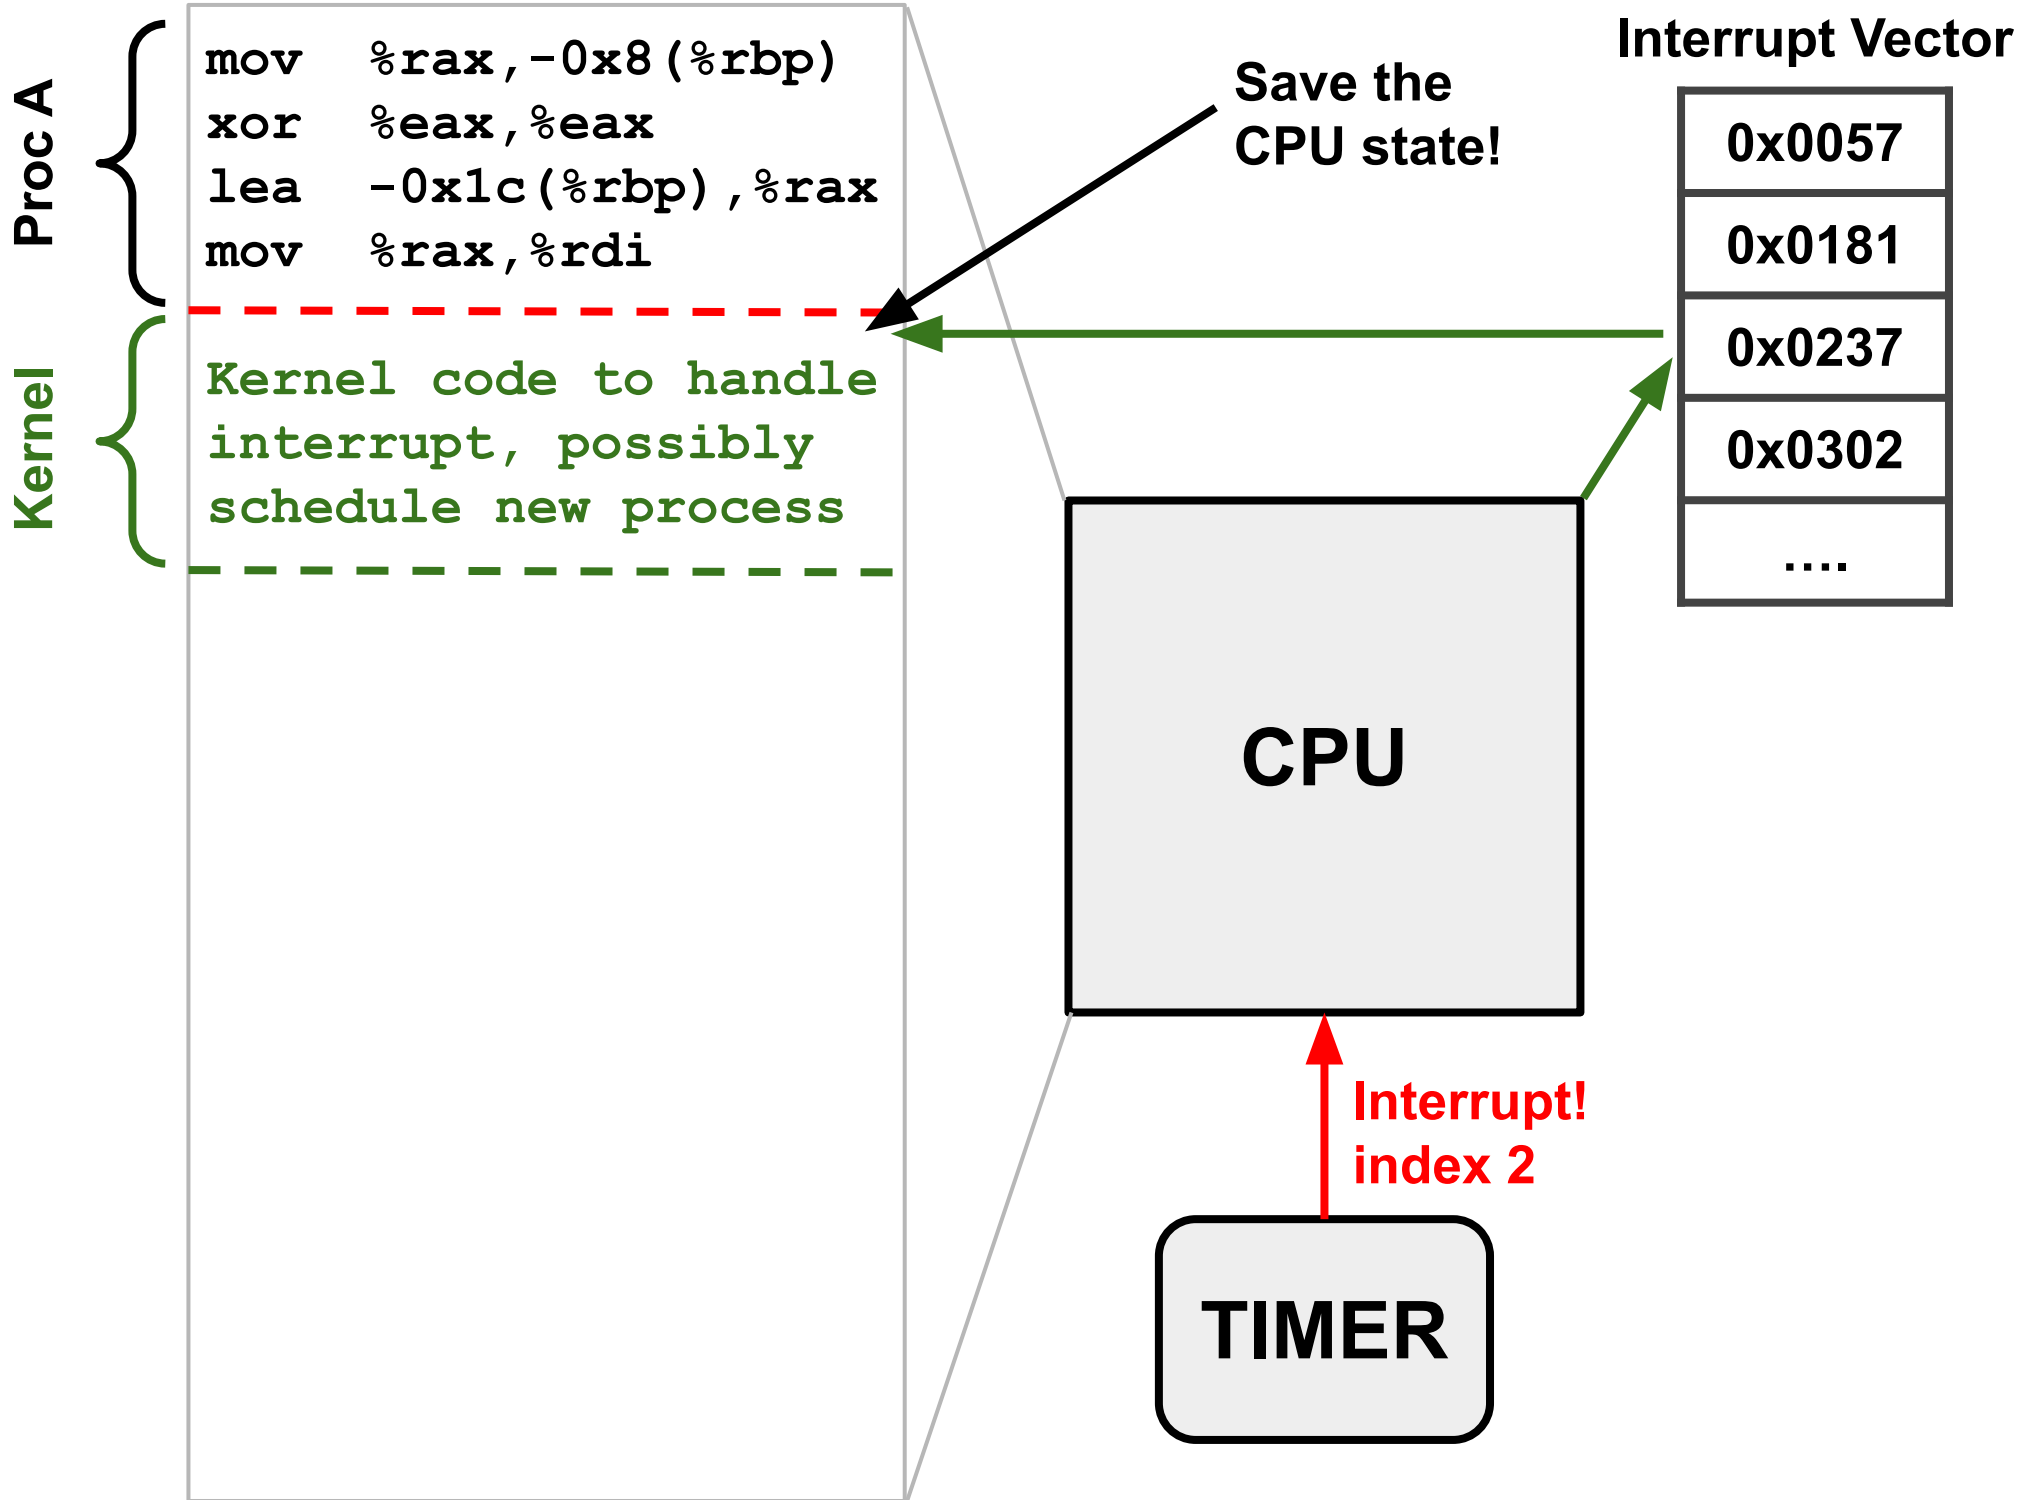
....

CPU

TIMER

Interrupt!
index 2





Proc A

```
mov    %rax, -0x8(%rbp)
xor     %eax, %eax
lea     -0x1c(%rbp), %rax
mov     %rax, %rdi
```

Kernel

```
Kernel code to handle
interrupt, possibly
schedule new process
```

Proc B

```
mov     %rax, %rcx
mov     $0x5, %edx
mov     $0x1, %esi
lea     0x0(%rip), %rdi
```

CPU

TIMER

Interrupt Vector



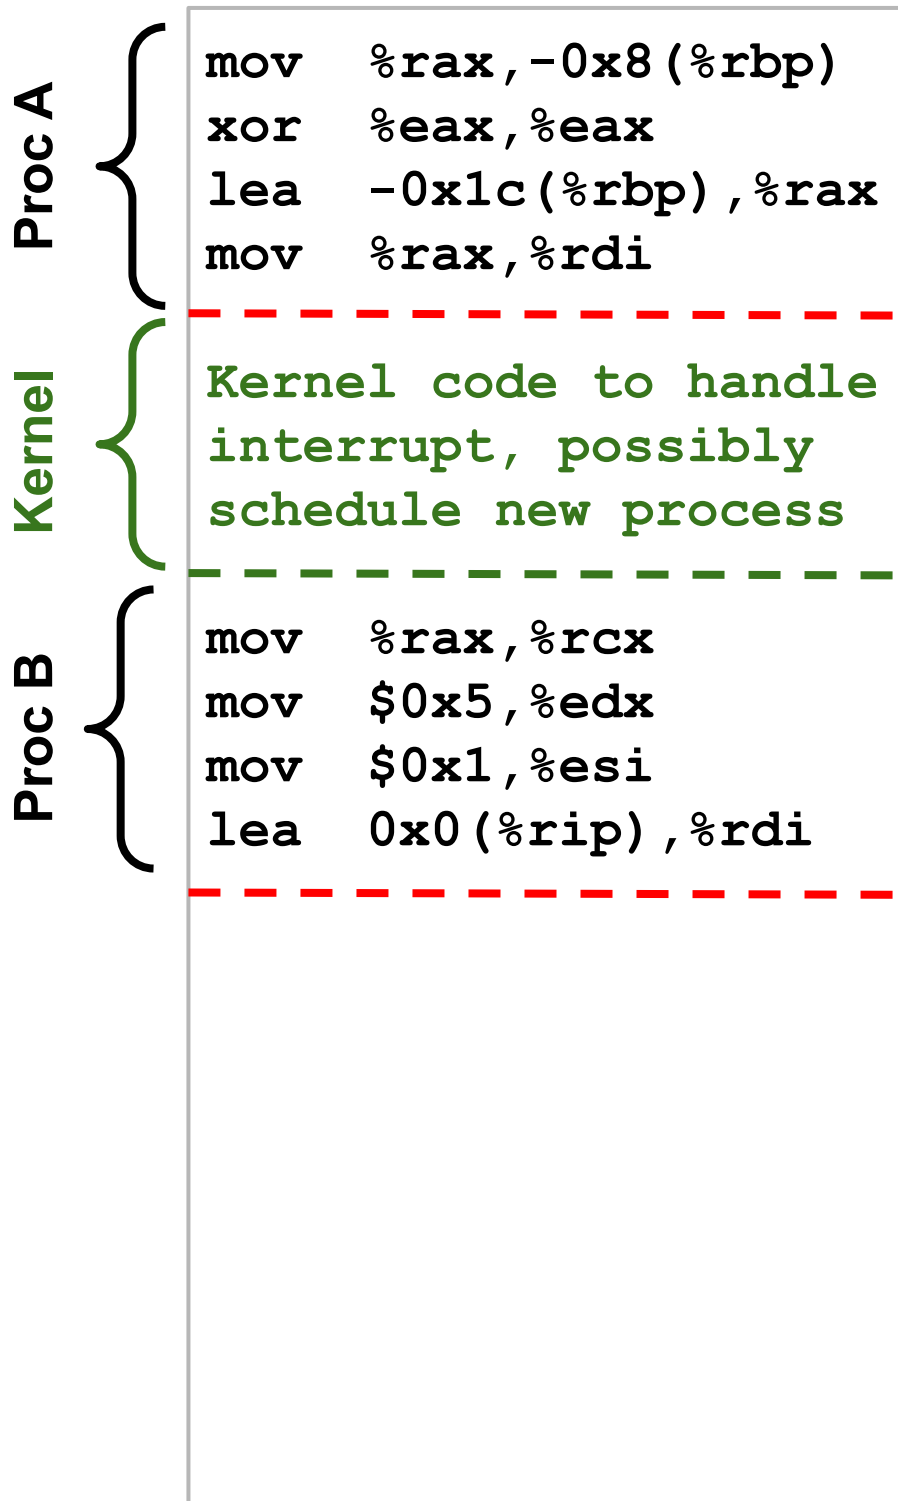
0x0057

0x0181

0x0237

0x0302

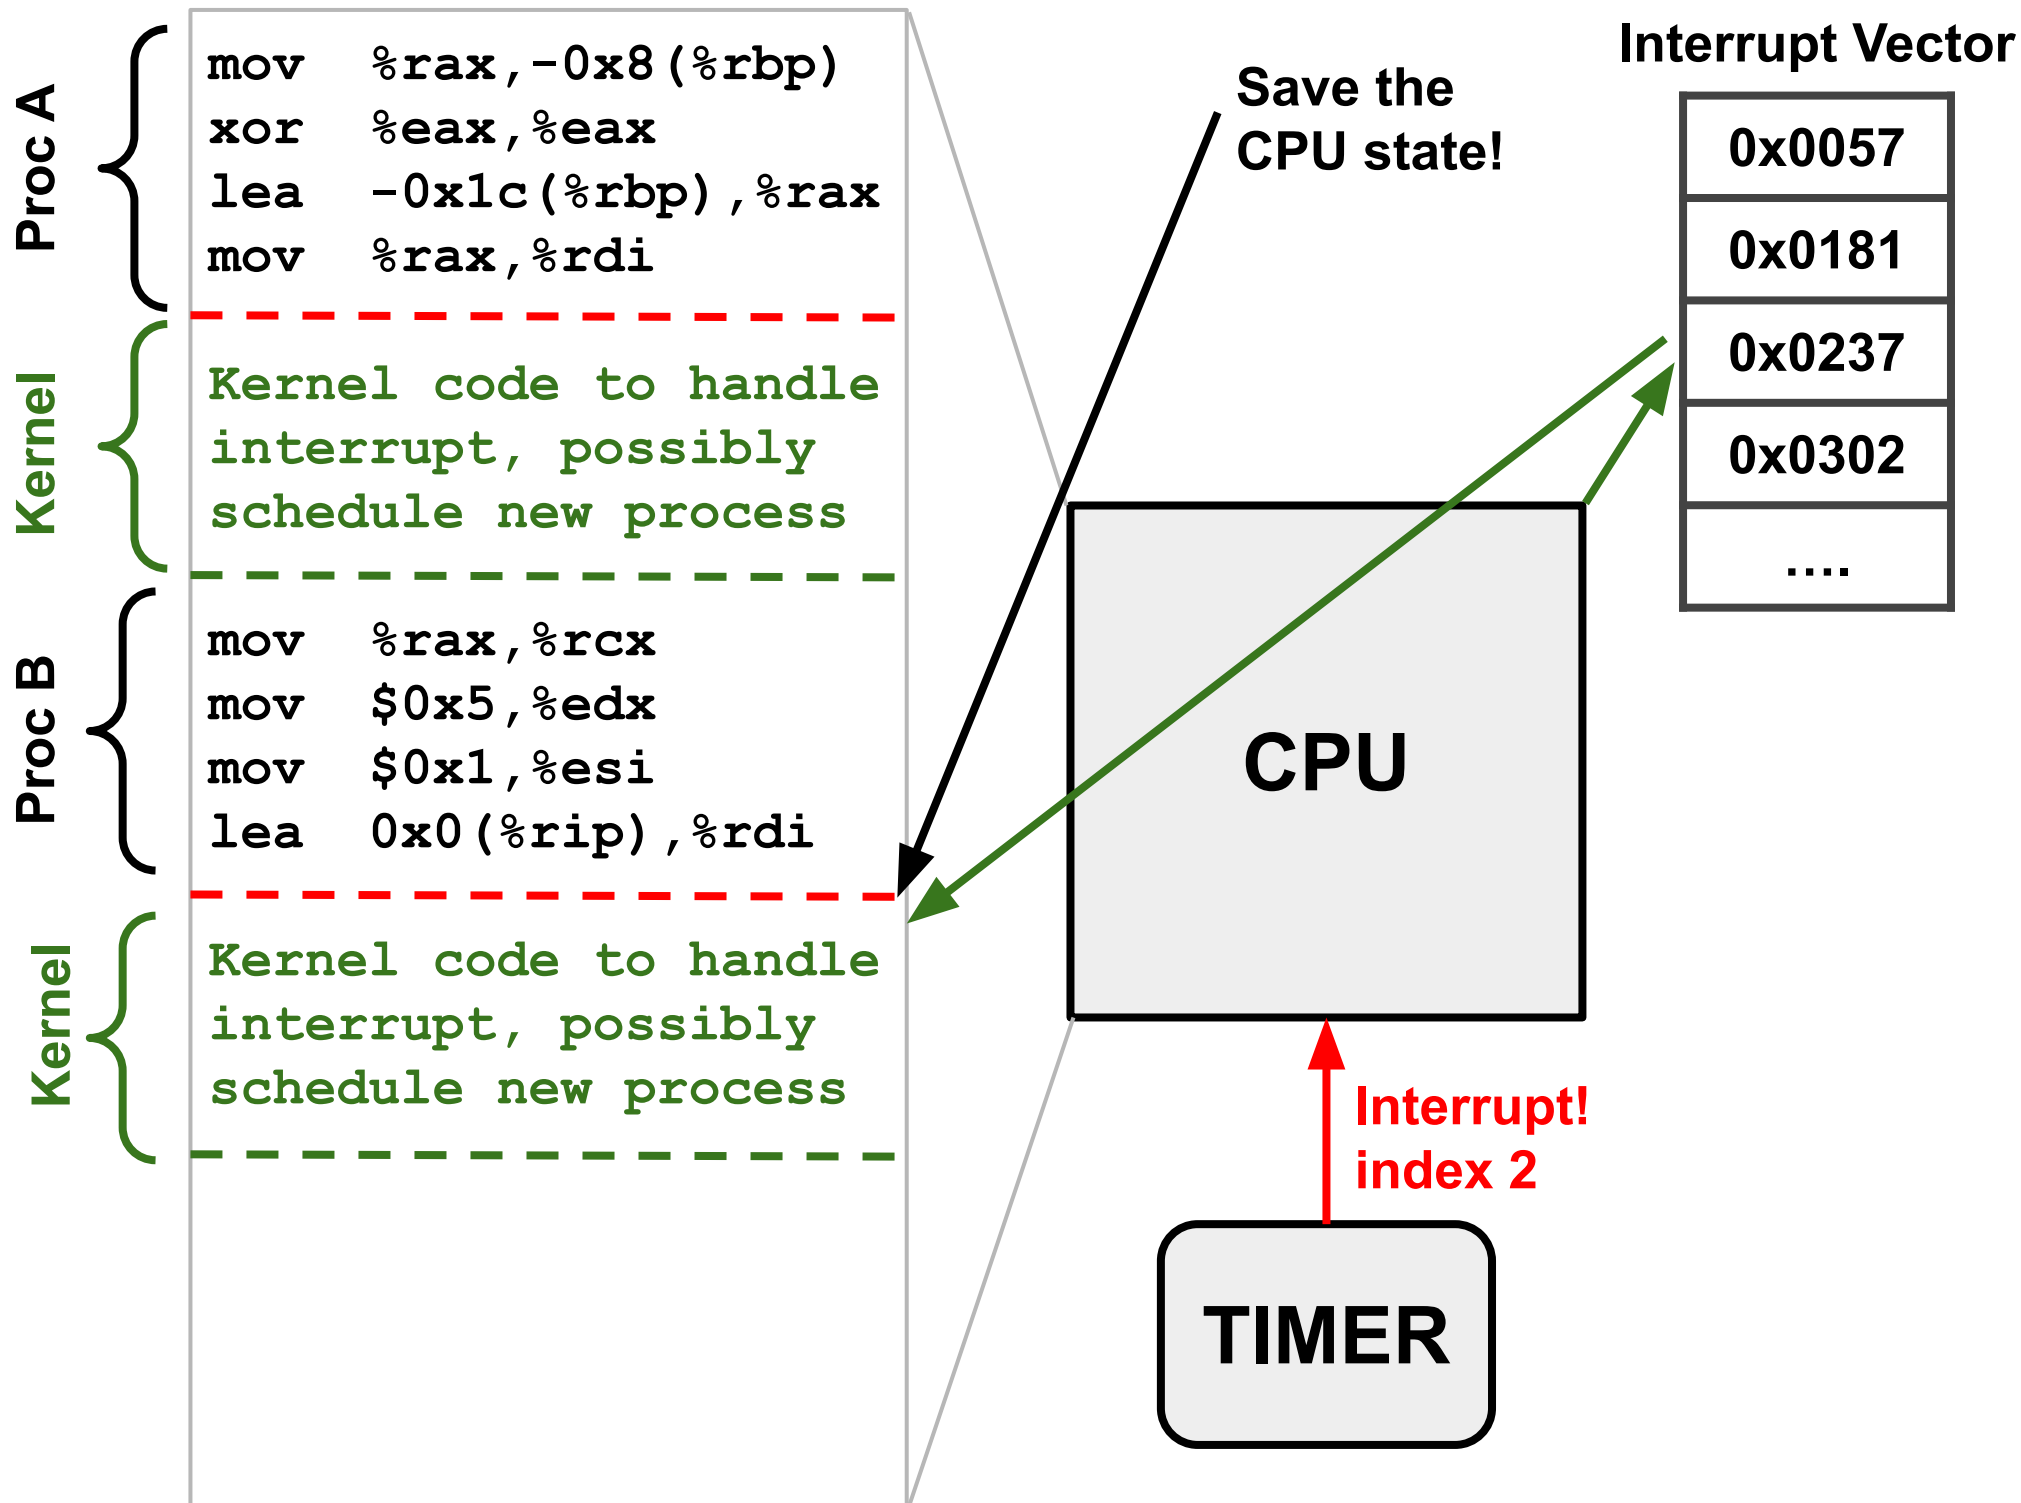
....



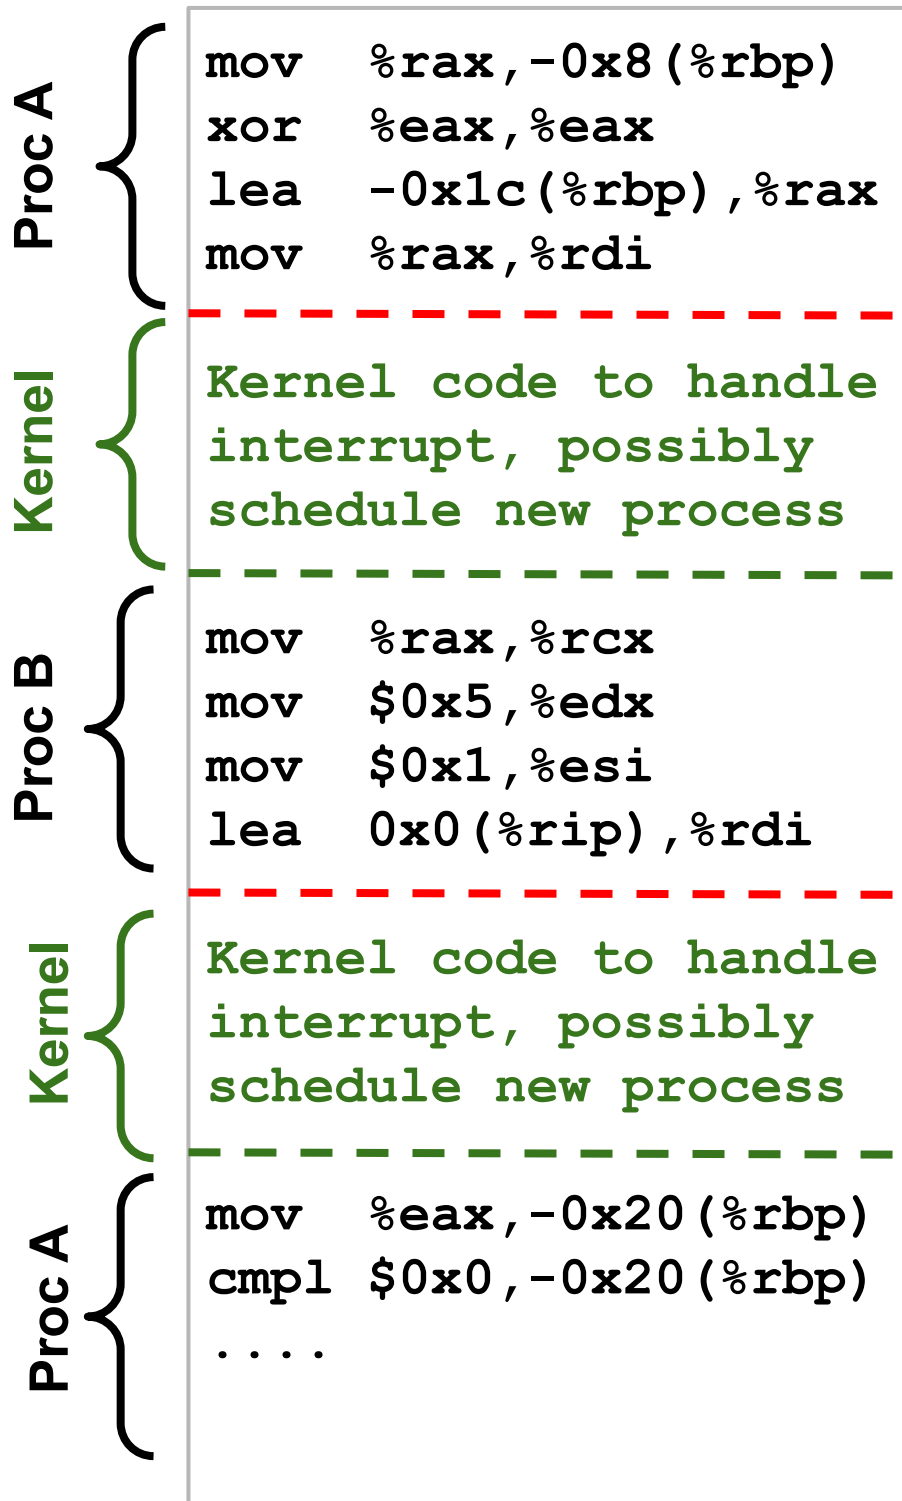


Interrupt!
index 2

Interrupt Vector

0x0057
0x0181
0x0237
0x0302
....





Interrupt Vector

0x0057
0x0181
0x0237
0x0302
....

Activity: Interrupts

Consider this Scenario:

A CPU runs at 3.0GHz. This is a simple CPU that executes exactly 1 instruction per clock cycle. The timer interrupt fires off 1000 times per second. How many instructions can be executed in each time slice?

Software and Device Interrupts

Visual Example

Proc A

```
mov    %rax,%rcx  
mov    $0x5,%edx  
mov    $0x1,%esi
```

CPU

TIMER

Storage
Device

Interrupt Vector

0x0057

0x0181

0x0237

0x0302

....

Proc A

```
mov    %rax,%rcx
mov    $0x5,%edx
mov    $0x1,%esi
int    0x2
```

CPU

TIMER

Storage
Device

Interrupt Vector

0x0057

0x0181

0x0237

0x0302

....

Proc A

```
mov    %rax,%rcx  
mov    $0x5,%edx  
mov    $0x1,%esi  
int    0x1
```

**Interrupt!
index 1**

Interrupt Vector

0x0057

0x0181

0x0237

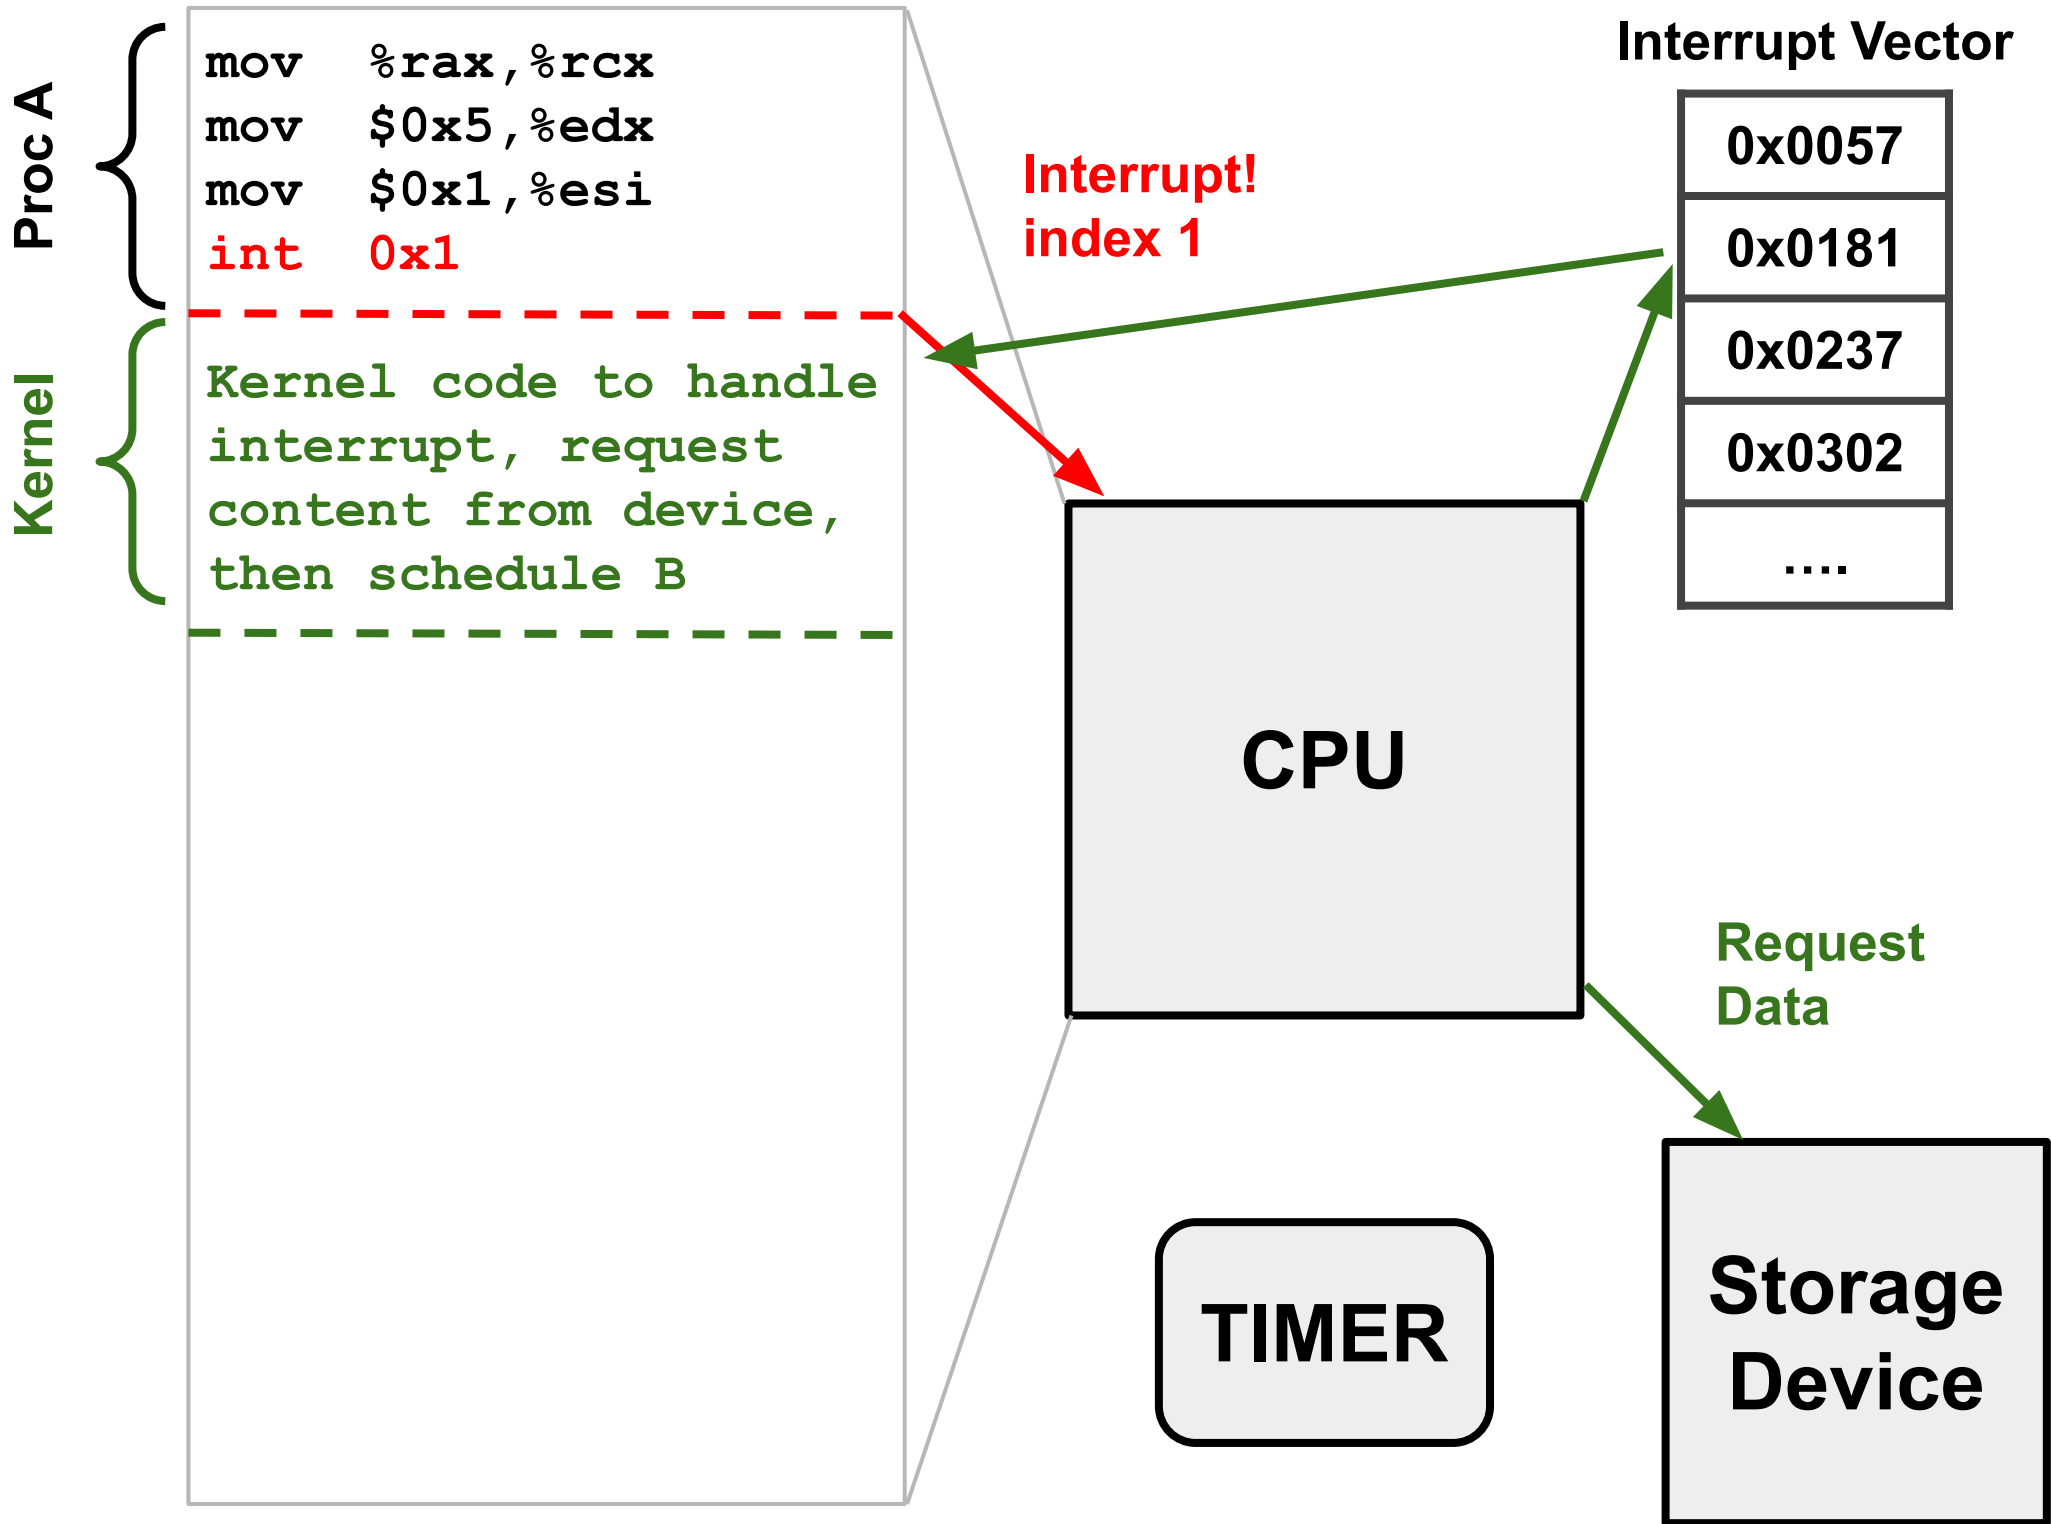
0x0302

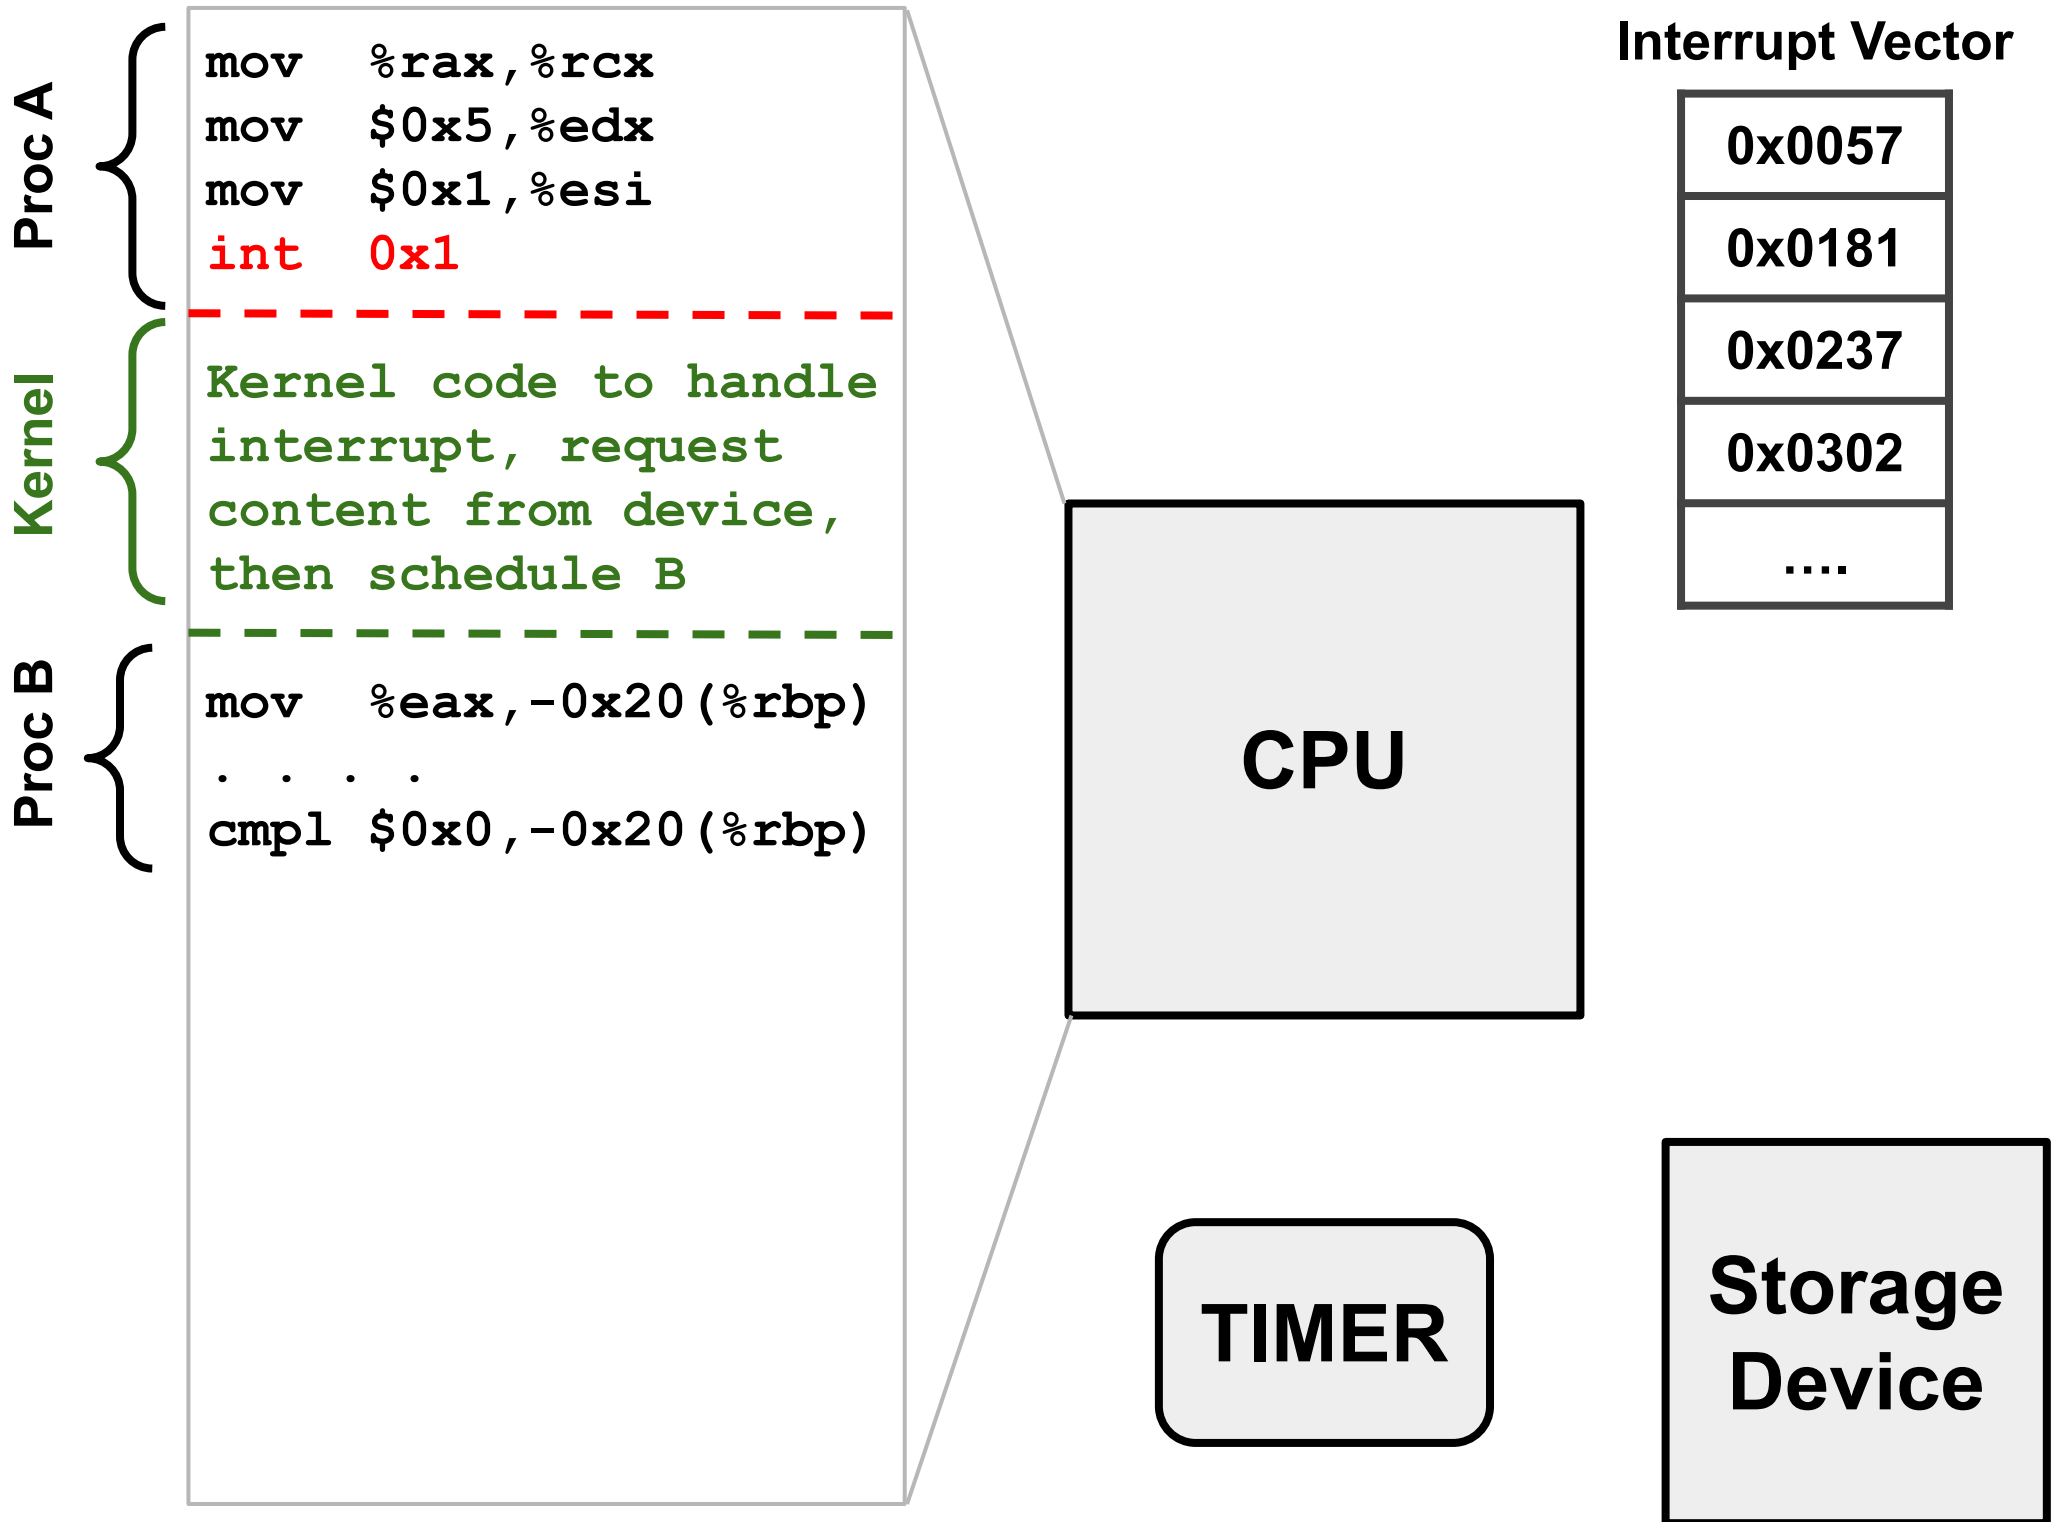
....

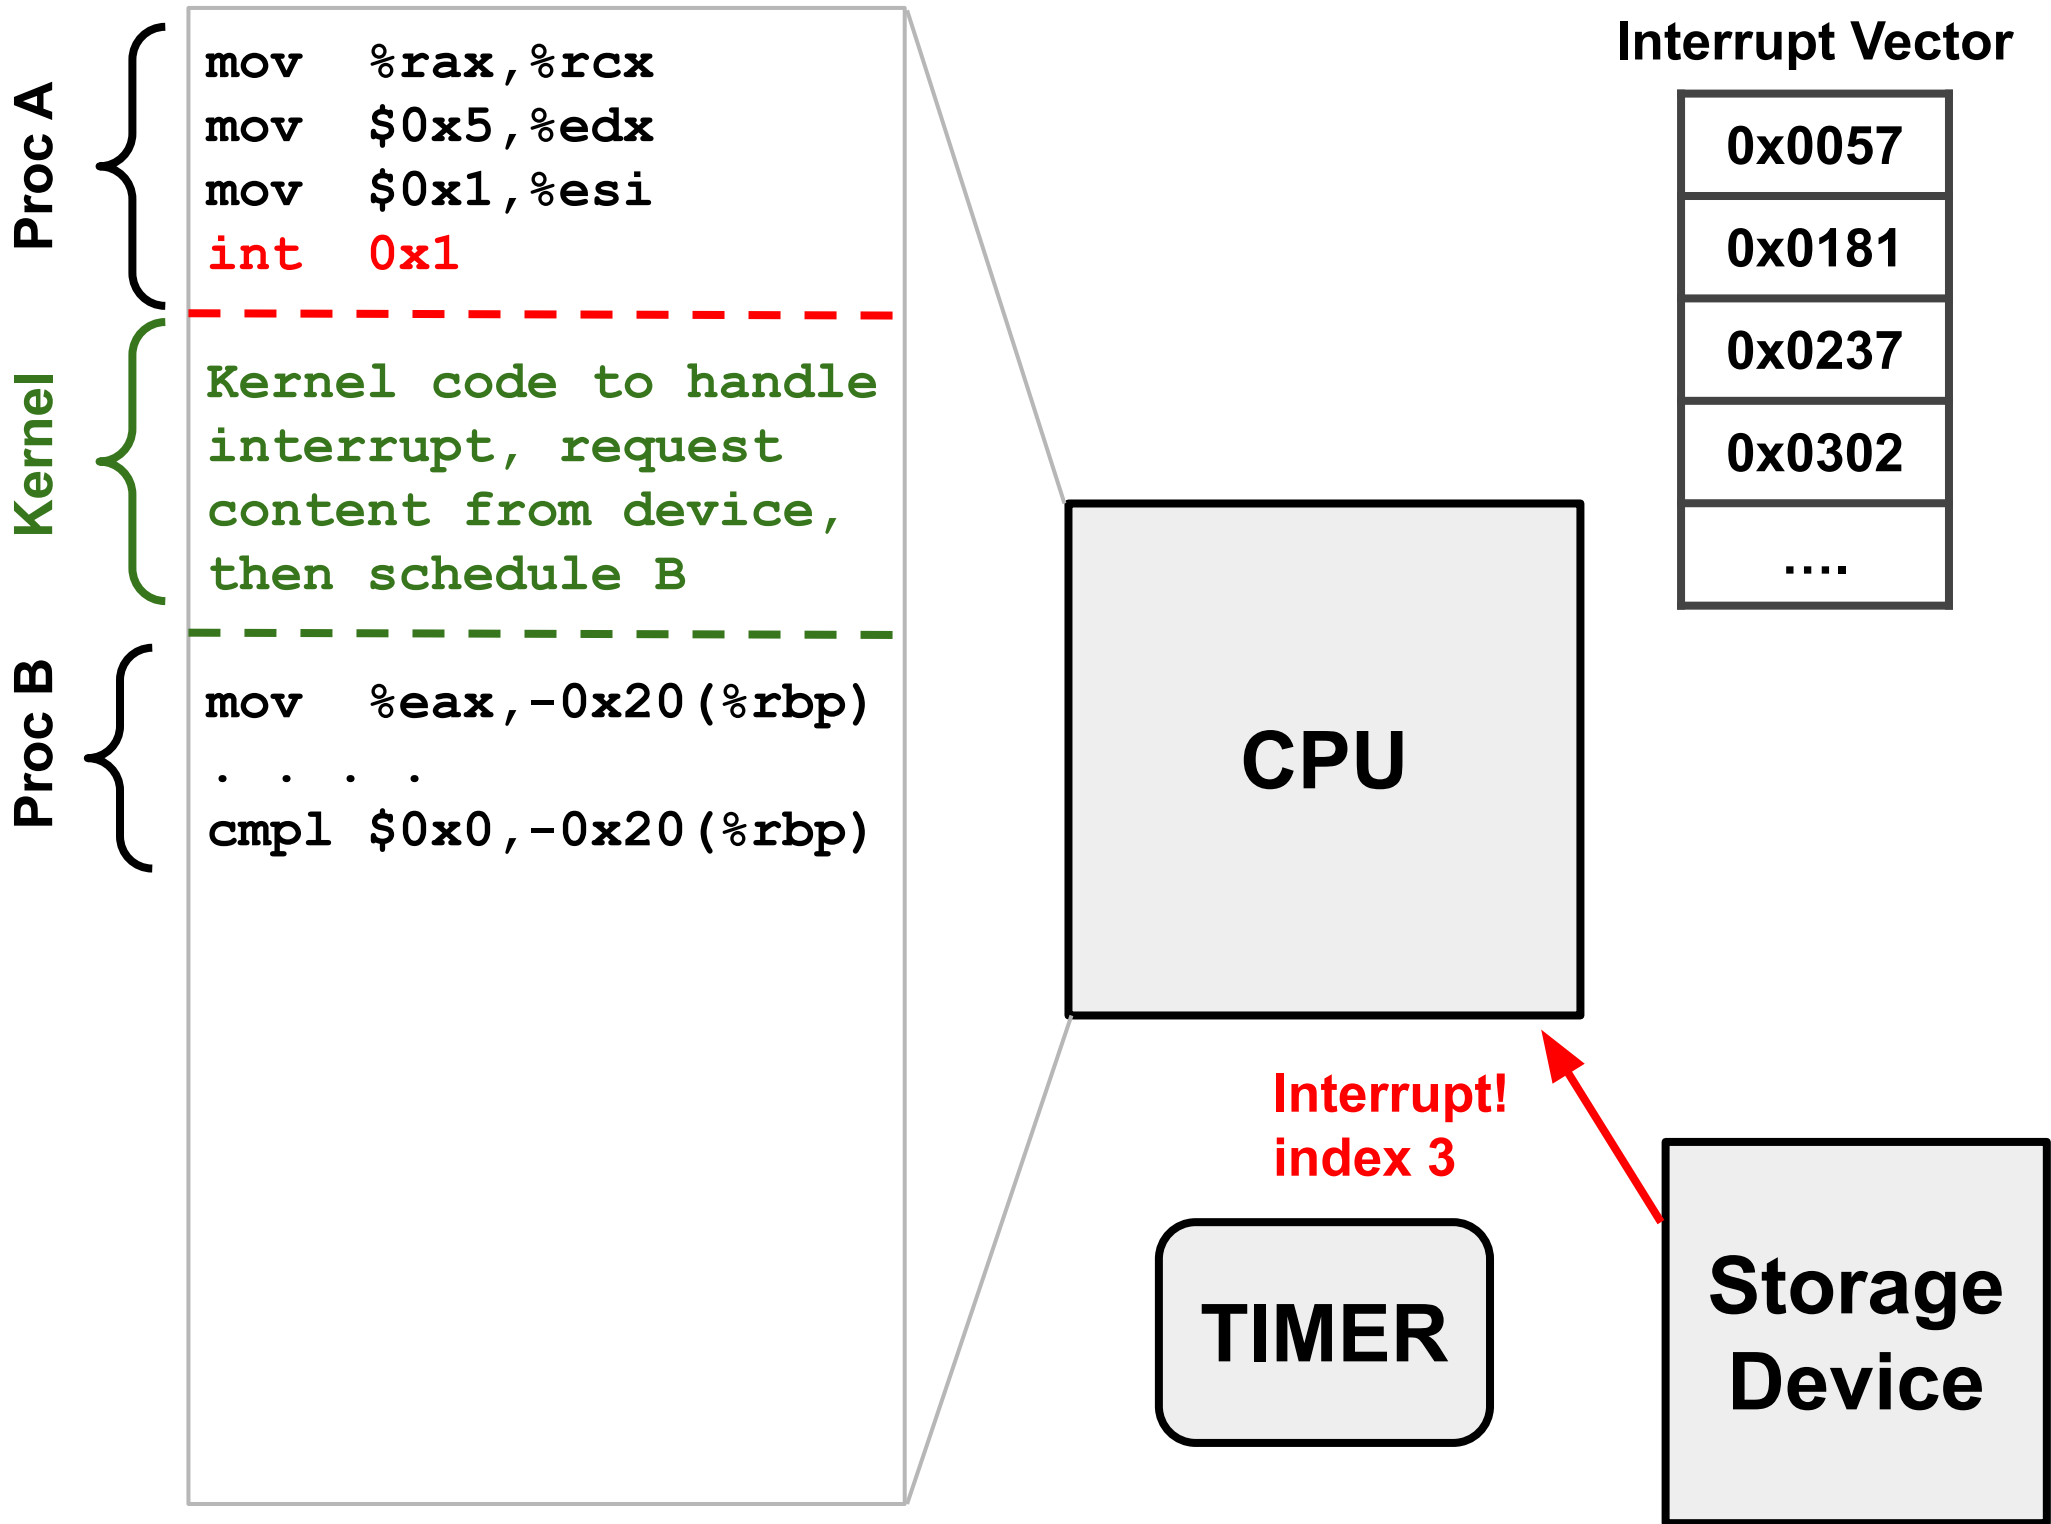
CPU

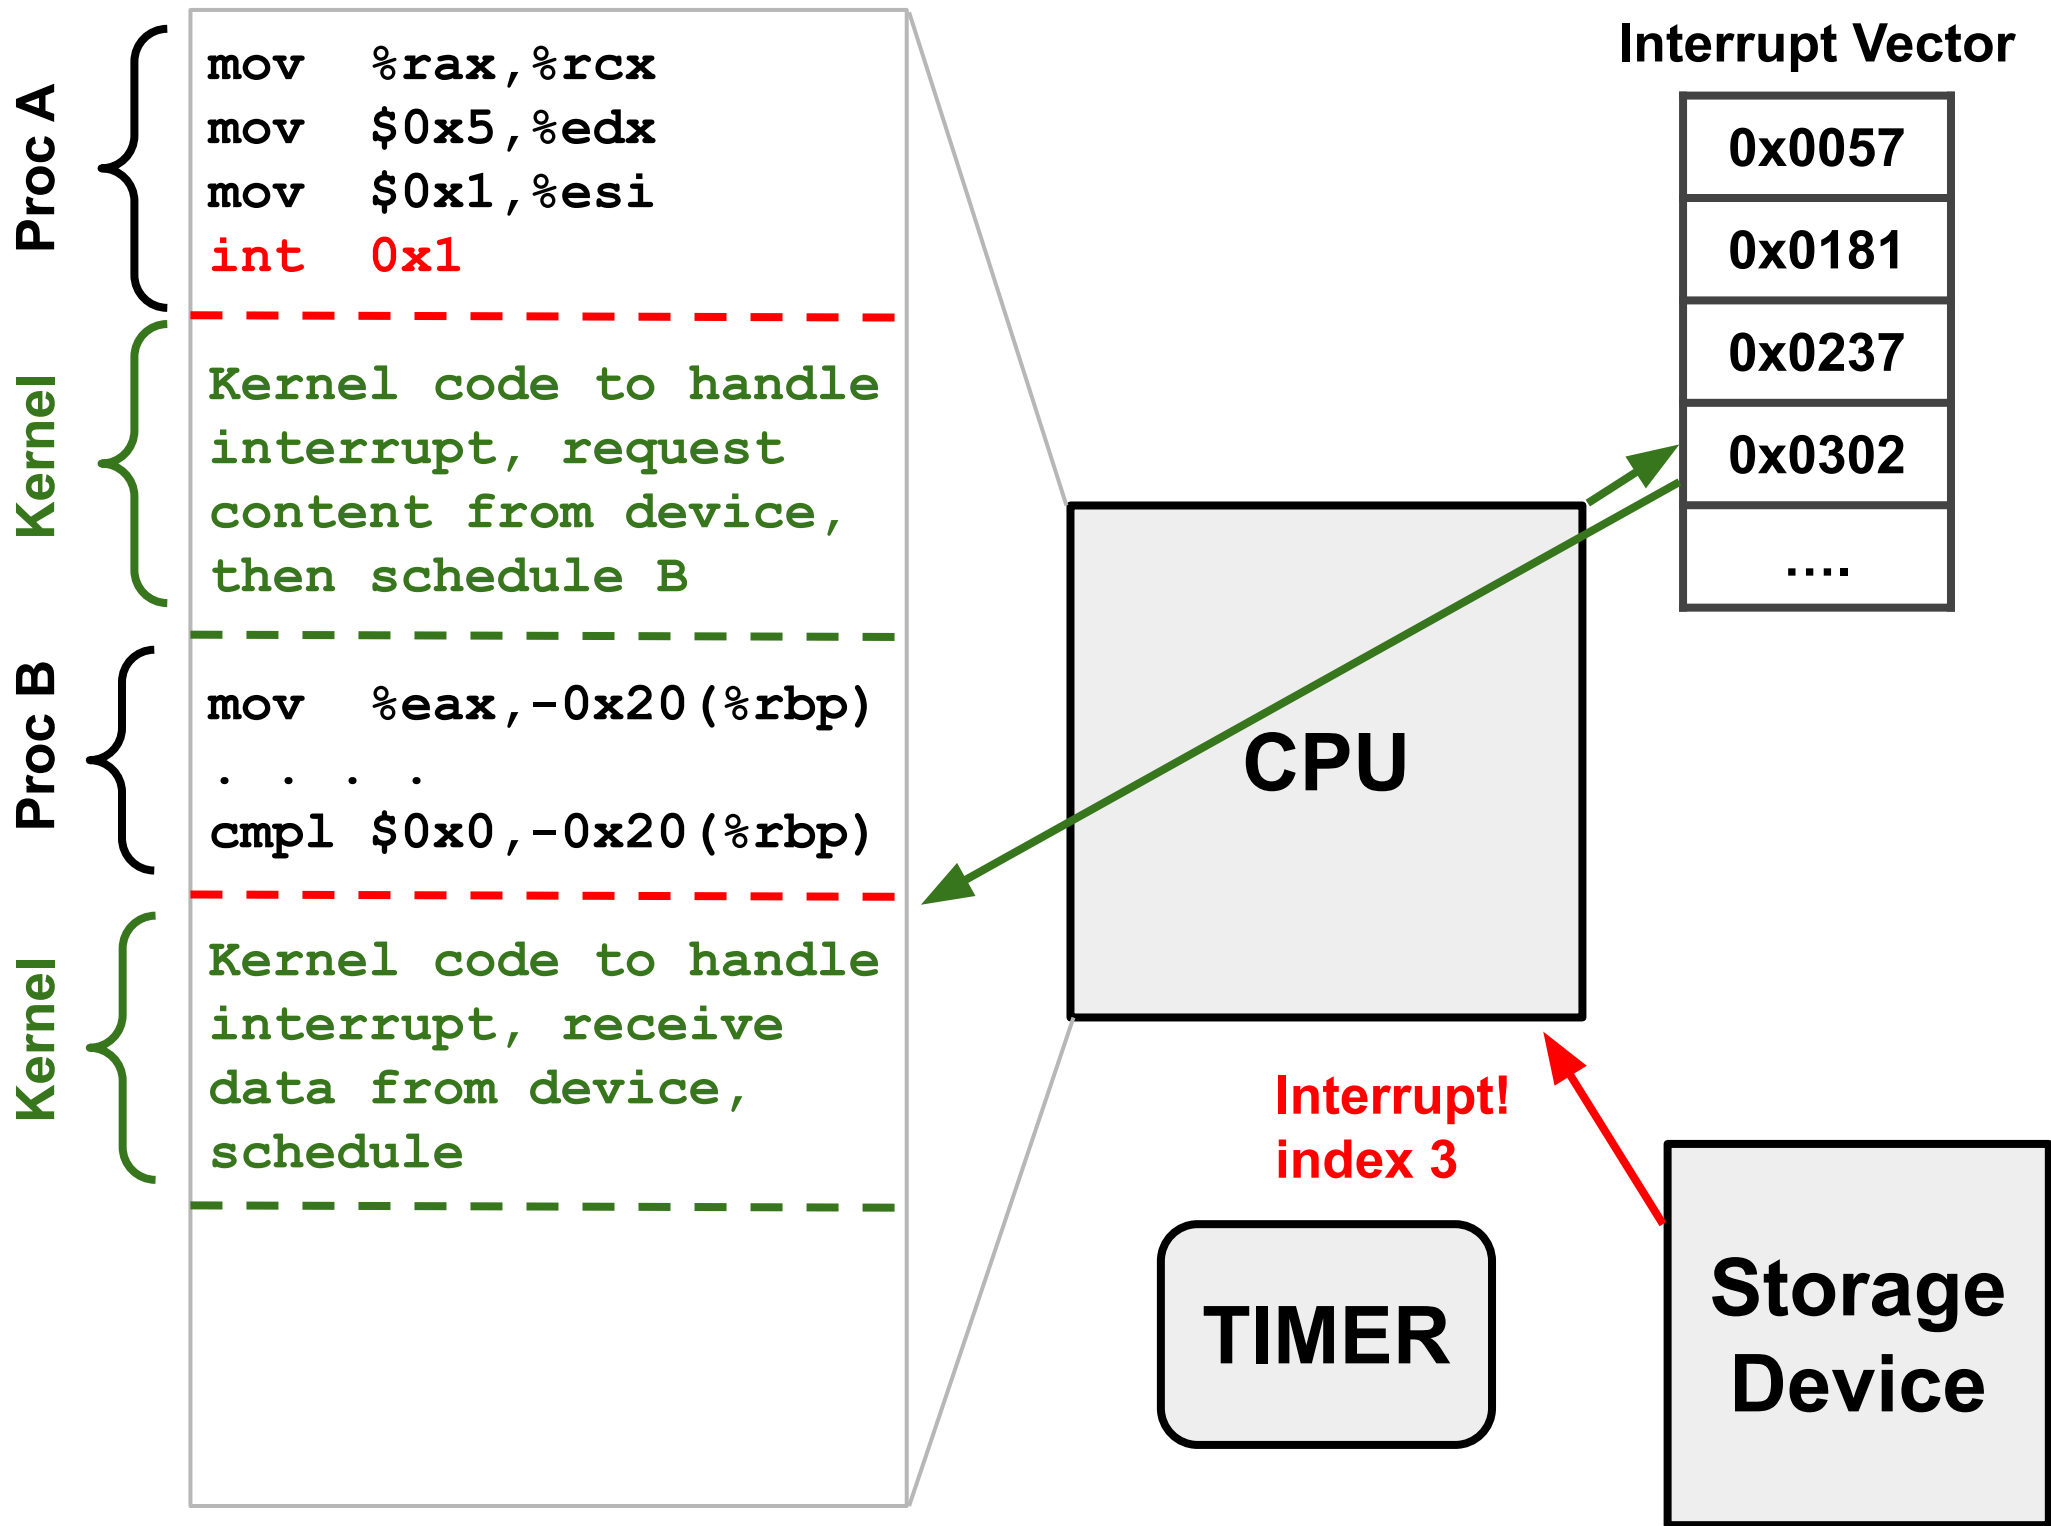
TIMER

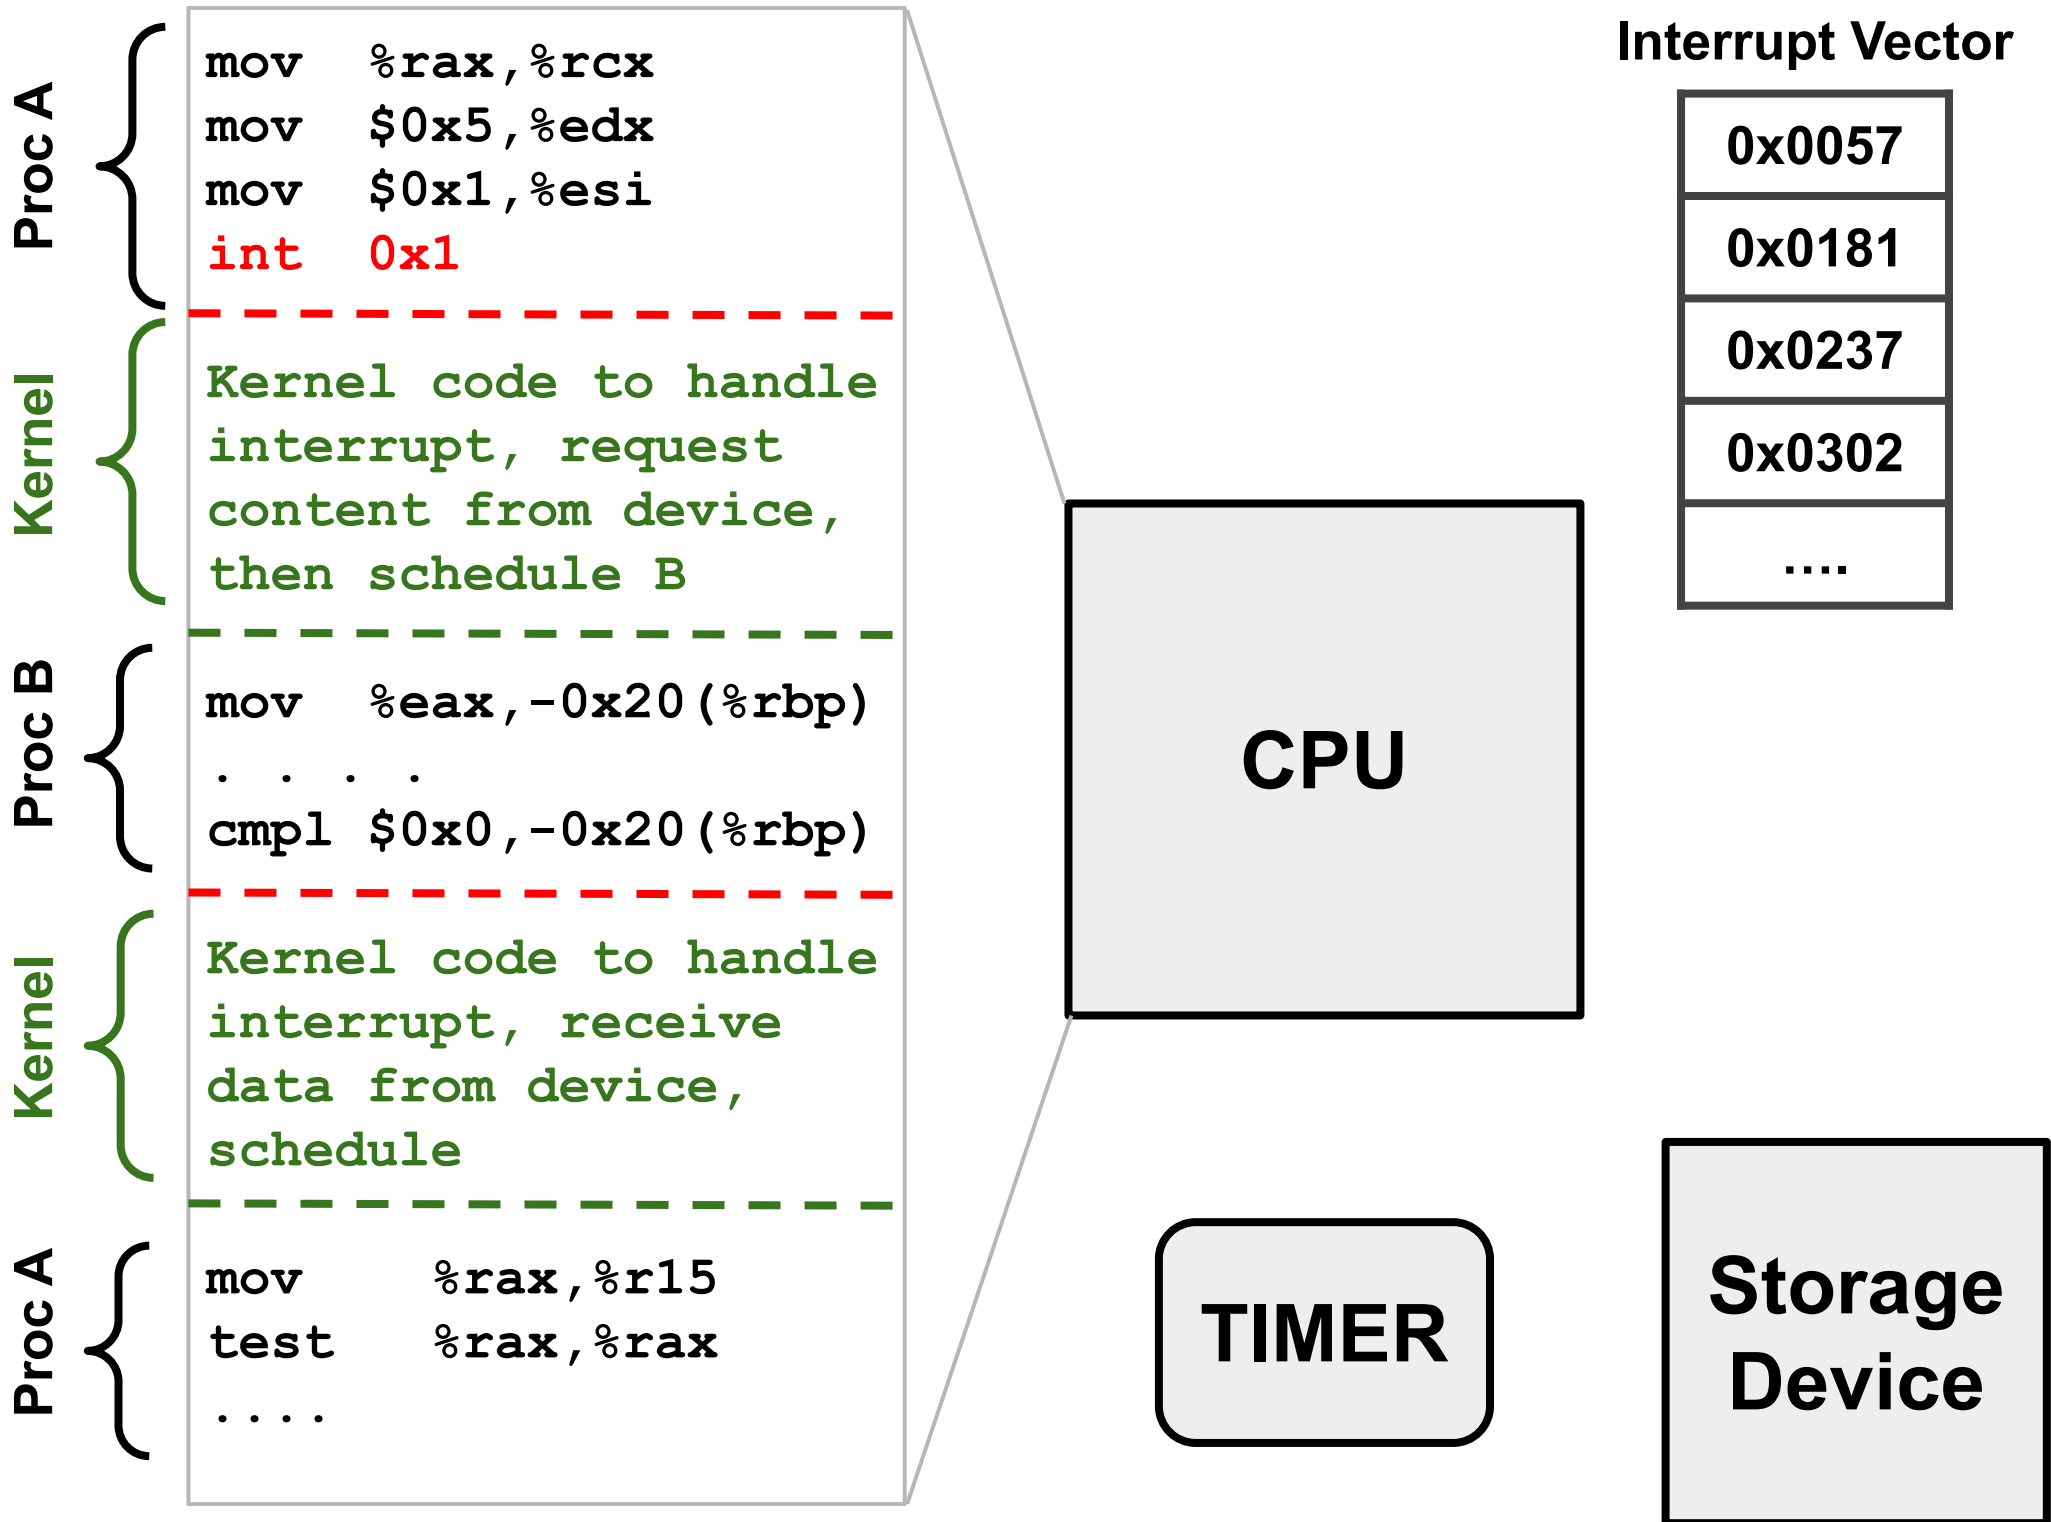
**Storage
Device**











Activity: Interrupts

Consider this scenario: Process A was running, then a timer interrupt occurs and the interrupt handling code starts executing

- What if the interrupt handler *modifies* a variable that process A cares about?
- What if the interrupt handler *uses* a variable that process A was changing?

Kernel A



```
.. instructions ..  
lea proc_count_addr %r1
```

CPU

TIMER

Interrupt Vector

0x0057

0x0181

0x0237

0x0302

....

Kernel A

```
.. instructions ..  
lea proc_count_addr %r1
```

Kernel

```
Kernel code to handle  
interrupt, decide if  
different proc should  
be scheduled
```

Interrupt Vector

0x0057

0x0181

0x0237

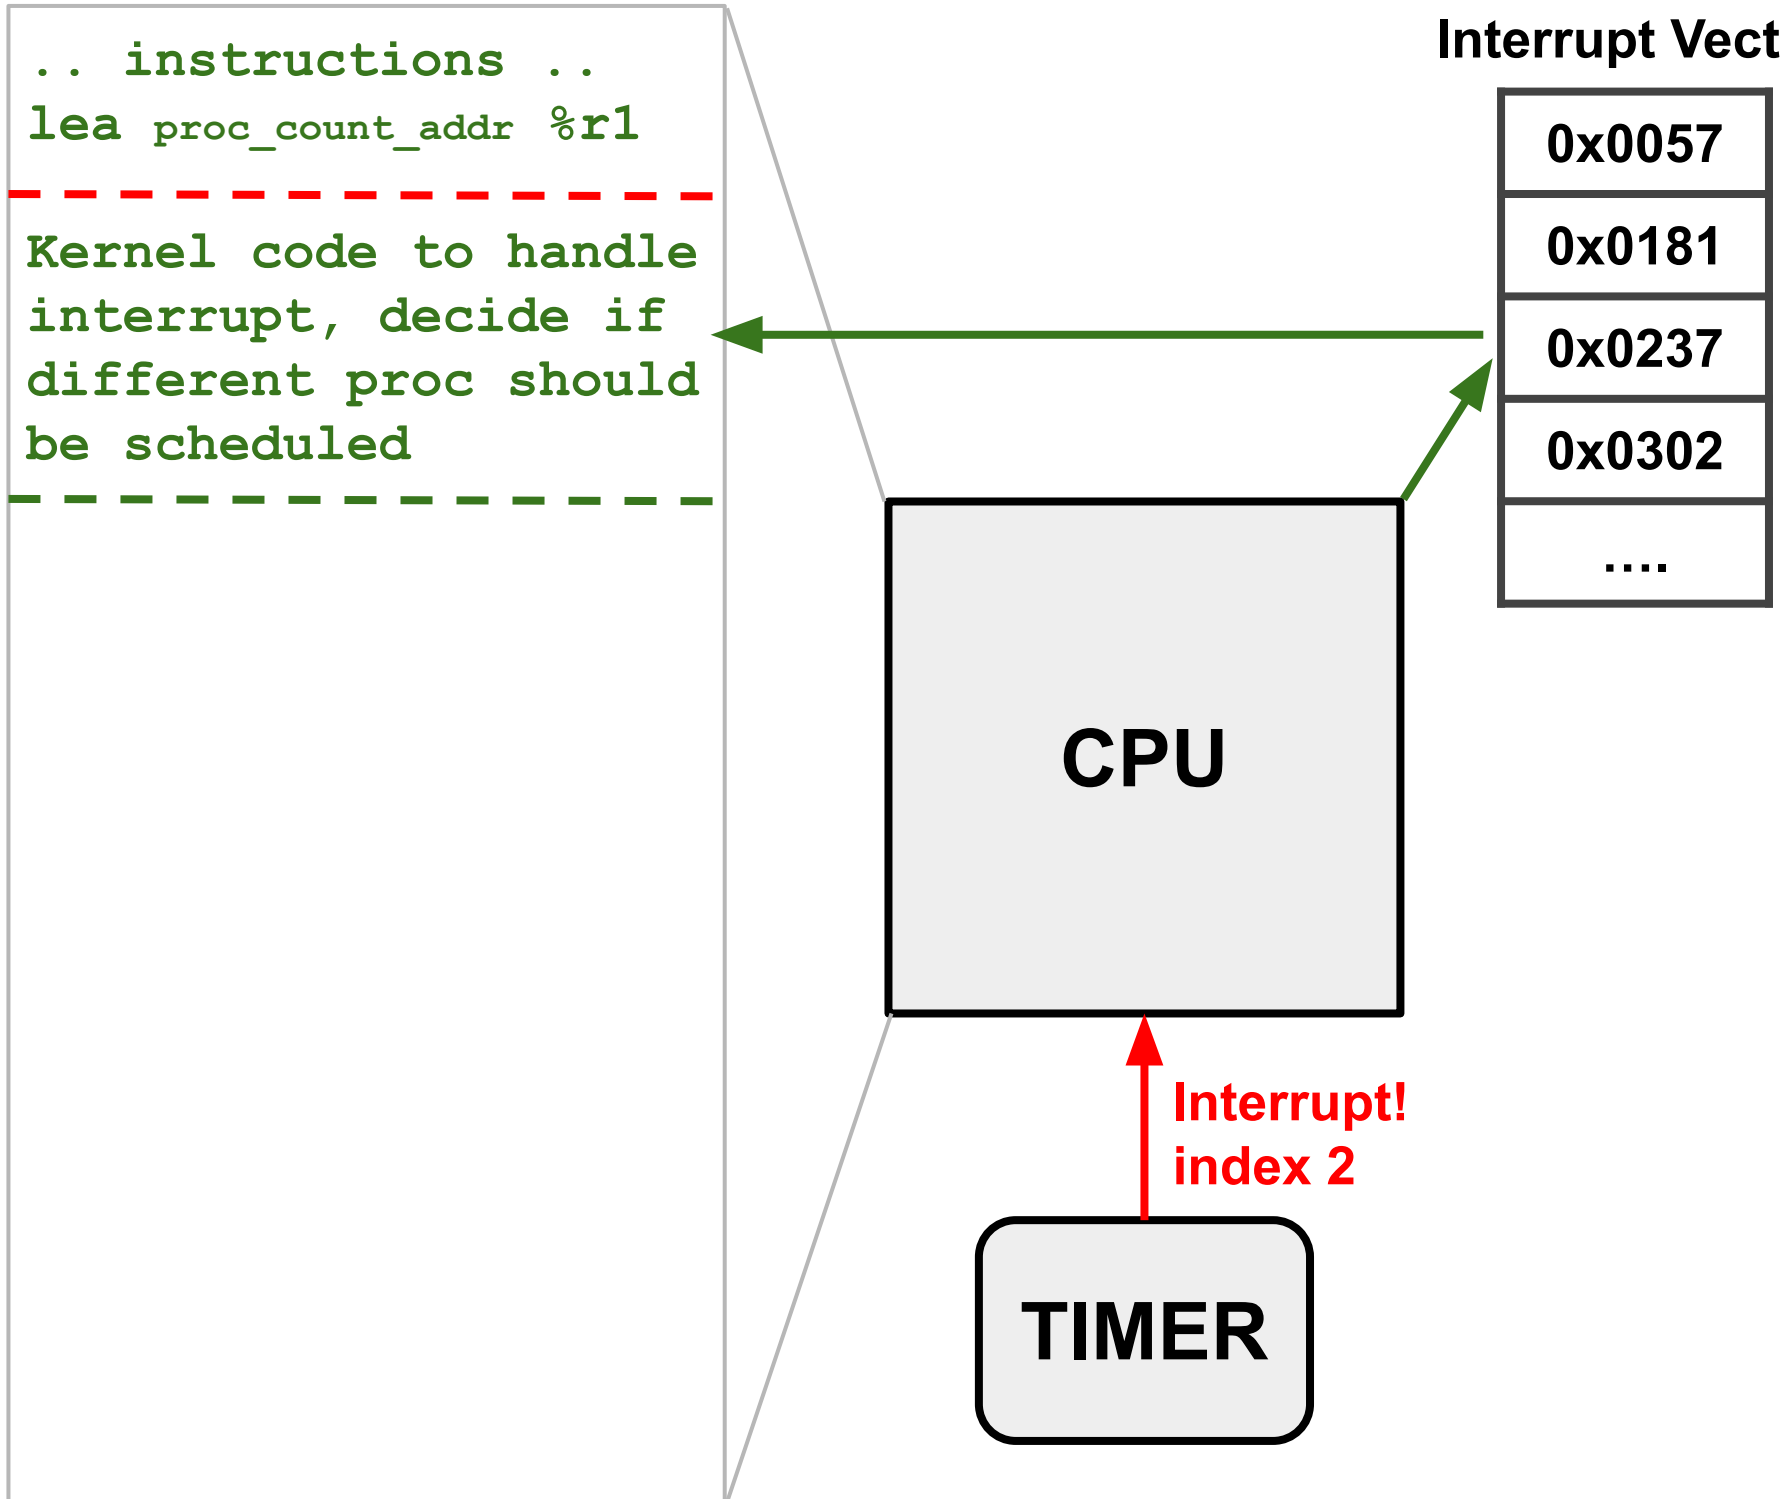
0x0302

....

CPU

Interrupt!
index 2

TIMER



Kernel A

```
.. instructions ..  
lea proc_count_addr %r1
```

Kernel

```
Kernel code to handle  
interrupt, decide if  
different proc should  
be scheduled
```

Proc X

```
mov  %rax,%rcx  
.  
.  
lea  0x0(%rip),%rdi
```

CPU

TIMER

Interrupt Vector

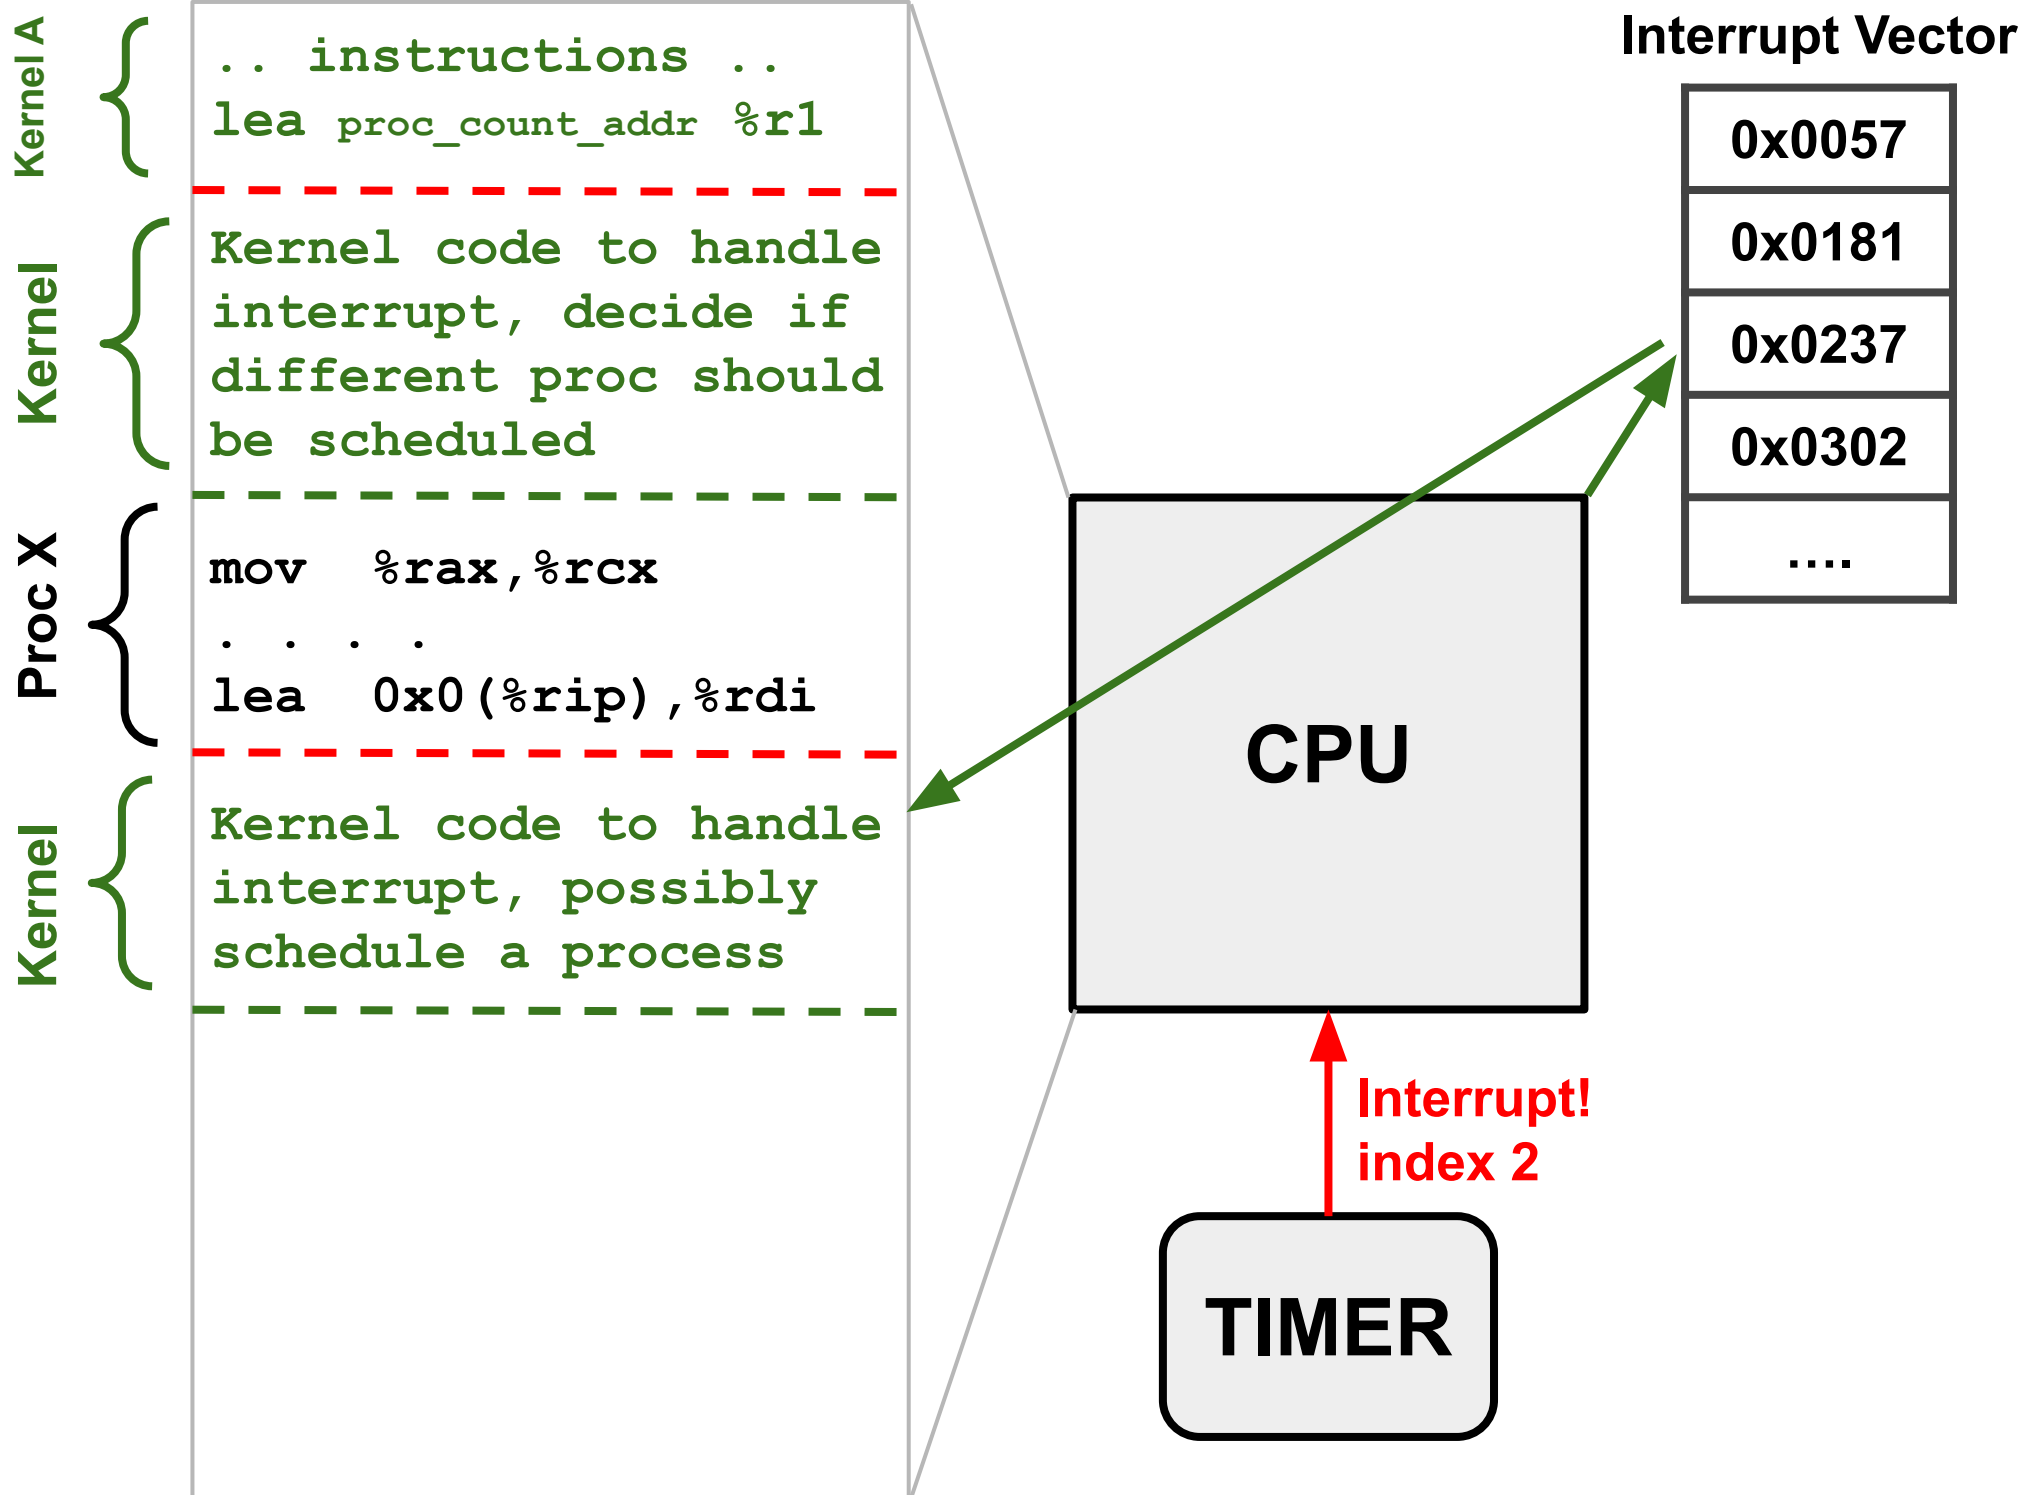
0x0057



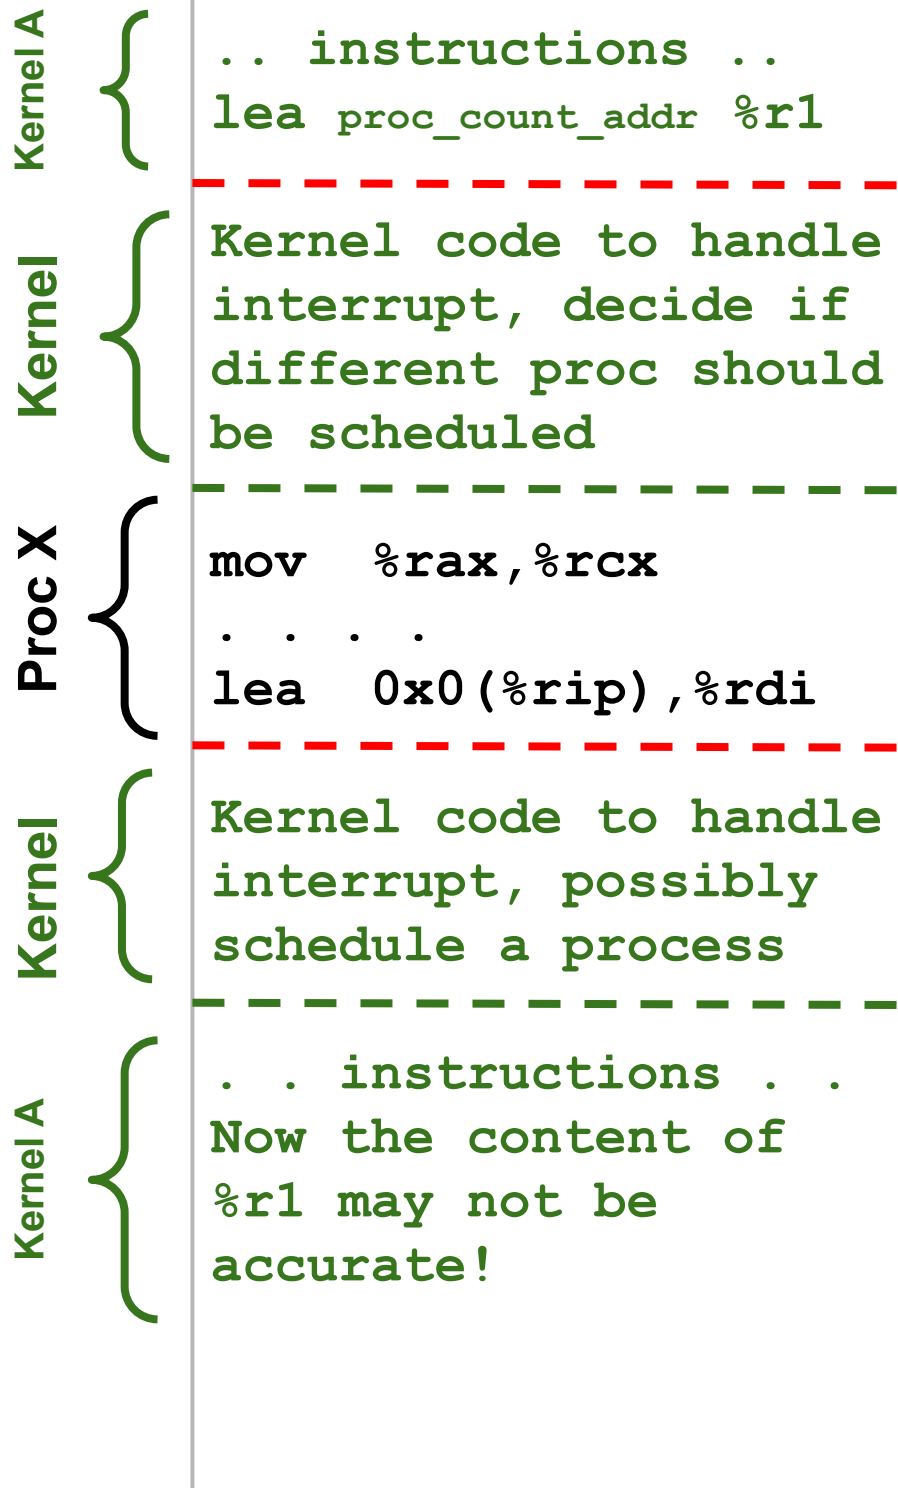
0x0181

0x0237

0x0302

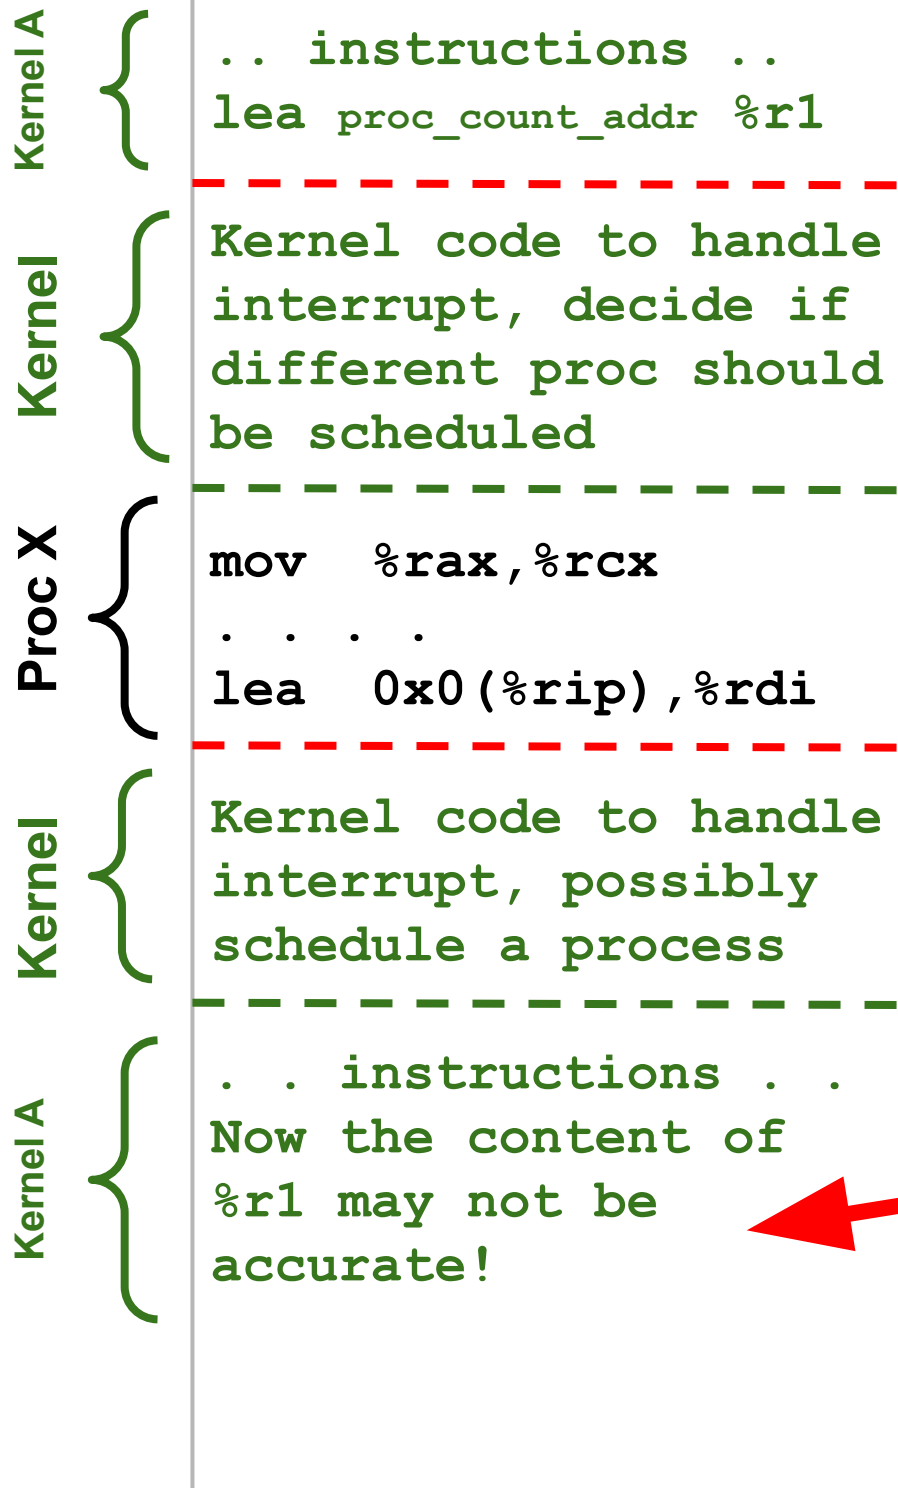
....







Interrupt Vector

0x0057
0x0181
0x0237
0x0302
....



Interrupt Vector

0x0057
0x0181
0x0237
0x0302
....



Problematic!



Races

- A **race condition** is when the outcome of a calculation depends on “accidental” (that is, unpredictable) details of how quickly it runs
- Remember: it is **impossible** to reliably predict your speed
 - Might be interrupted
 - Might be context-switched
 - Cache, paging, etc.

Races

- Races are **really bad!**
- Often will seem to work, but fail randomly
- Very difficult to replicate
- Very difficult to test your fix
- Conclusion: **prevent them before they happen!!!**

Interrupts

- Kernel code (not user!) can **disable interrupts** at any time
 - No interrupts will fire
 - No interrupt handlers will run
 - External interrupts still happen, CPU remembers them
 - Interrupts fire **immediately** when user re-enables interrupts
 - Some interrupts cannot be disabled

Interrupts

- In Phase 1 you will need to ensure interrupts get disabled whenever a transition to the kernel happens.
- Set the ***Processor Status Register*** (PSR) to the appropriate value.

Activity: Interrupts

- Disabling interrupts works well to solve these problems on a ***single-core*** CPU system
- What about multi-core or multi-cpu systems?

Preventing Concurrency

- In a **single-CPU** OS is simpler! If you disable interrupts:
 - Time-slicing never happens
 - Interrupts can't run
- In a multi-CPU OS, concurrency is **real**
 - Use locks to protect data
 - When self-deadlock is a worry, disable interrupts **before** you gain a lock

Concurrency

Consider this code

```
while True:
    load  x → $r1
    inc   $r1
    store $r1 → x
```

Critical Sections

A **critical section** is a portion of the program where interrupting it might cause a race

```
while True: } Non-Critical
    load  x → $r1
    inc   $r1
    store $r1 → x } Critical
                  Section
```


Locks

- A **spinlock** is a lock where `gain()` is implemented as a tight `while` loop
 - In some implementations, CPU can be stuck forever
 - In others, the process eventually gives up and goes to sleep

Locks

- A lock can only have one owner
- To become the owner, you “gain” (or “lock”) it
 - Will block if it is already owned
- To release ownership, you “release” (or “unlock”) it

Locks

- **WARNING**
- Locks are only useful if you use them in all the right places. Locks don't truly protect data; they just block processes from reaching code.
- If you forget to use `gain()` / `release()` on one CS, it will be a *danger* to all the other CSes (and vice-versa).

Interrupts

- **Self-deadlock** is the condition when a process owns a given lock, but is also blocked, trying to gain the lock a second time
- Because the lock will never be released, the lock will never be gained
- Thus, we're stuck forever
- Interrupts can trigger self-deadlock
 - (Other things can cause it, too...)

Interrupts

- What are the tradeoffs of disabling interrupts?

Good: self-deadlock impossible

Bad: preemptive context switches never happen

Preventing Concurrency

- Instead of gaining & releasing any locks, we will simply enable and disable interrupts!

- while True: # not in CS
 old_psr = disable_ints()
 load x → \$r1 # in CS
 inc \$r1 # in CS
 store \$r1 → x # in CS
 restore_ints(old_psr)

Concurrency

Three Classic forms of Concurrency

Multiprocessing

We won't be doing this in USLOSS

Time-Sharing

Will need to be considered

Interrupts

Exist in USLOSS!