

Messages

Messages for Communication & Concurrency Control

IPC

Our Message Protocol

Mutexes & Signaling

Messages

IPC

IPC

- **Interprocess communication** (IPC): A system that allows processes to talk to each other

IPC

- **Interprocess communication** (IPC): A system that allows processes to talk to each other

What are some useful applications of IPC?

IPC

- IPC can denote communication between processes on the same computer
- However, IPC can also be used for **distributed processes**
 - Processes run on different computers
 - Connected through network
 - Sockets or messaging most common

IPC

- **Interprocess communication** (IPC): A system that allows processes to talk to each other
 - Messages/mailboxes (phase 2)
 - Signals
 - Pipes, files
 - Semaphores (phase 3)
 - Shared memory

IPC

- **Mailboxes** allow processes to send **messages** to other processes
 - Atomic messages, small payload
 - Buffering (often)
 - Blocking vs Non-blocking

IPC

- A **signal** is a software interrupt, sent from one process to another
 - OS must deliver, forces receiver into a **signal handler** function
 - Processes can block signals temporarily (analogous to disabling interrupts)
 - Typically no payload
 - “Default” behavior if no handler
 - Ignore, terminate, core dump, etc.

See `$ man 7 signal`

IPC

Why would one choose to use a message/mailbox over a signal for IPC?

IPC

- A **pipe** is a file-like mechanism which handles a stream of bytes, but is never stored on disk
 - `stdin, stdout, stderr`
 - `cmd1 | cmd2`
 - Usually only one reader, but many writers OK
 - Block writers when buffer fills

IPC

Example!
`/tmp/search.c`

- A **pipe** is a file-like mechanism which handles a stream of bytes, but is never stored on disk
 - `stdin, stdout, stderr`
 - `cmd1 | cmd2`
 - Usually only one reader, but many writers OK
 - Block writers when buffer fills

IPC

- **Semaphores** are integers used for counting events, block processes when insufficient resources
 - No payload
 - $P()$ decrements, blocks if insufficient
 - $V()$ increments, releases blocked procs
 - Useful for mutexes, resource control, buffering
 - Slide deck later!

IPC

- **Shared memory** is any time that two processes map the same writable virtual page
 - Probably map to different virtual addresses
 - Instantaneous data flow
 - Need additional mechanism if want to coordinate via blocking
 - Hard to dynamically resize
 - Example: `critical.c` and `critical_lock.c`

Activity: IPC

Two processes A and B need to achieve mutual exclusion in a critical section via a lock (mutex).

Which form of IPC should be used?

Activity: IPC

Two processes A and B need to both send and receive strings between each-other to coordinate a computation.

Which form of IPC should be used?

IPC

- **Non-blocking** functions allow a program to ask if something is available, but refuse to block if not
 - Mostly used for IPC & dist. communication
 - Great for handling multiple simultaneous connections
 - Often, confusing to use
 - Can hurt performance if you don't have some way to sleep (spinning or polling)

Messages

Phase 2 Messaging Protocol

Our Messaging Protocol

- In our simulation (Phase 2), we will implement mailboxes and messages
- Mailboxes are created with `MboxCreate()`, and identified by integer IDs
- Each mailbox will have two limits:
 - Max bytes per message (can be 0)
 - Max buffered messages (can be 0)

Our Messaging Protocol

- In our simulation (Phase 2), we will implement mailboxes and messages
- Mailboxes are created with `MboxCreate()`, and identified by integer IDs
- Each mailbox will have two limits:
 - Max bytes per message (can be 0)
 - Max buffered messages (can be 0)

Why would we ever want either or both of these to be zero?

Our Messaging Protocol

- `MboxSend()` will “buffer” a message when there is no receiver ready to receive it
 - Each mailbox has a max (can be 0)
 - Sender will block if can't buffer
- Buffered messages must be delivered in order!!
- Blocked senders must be kept in order, too
 - Must buffer in same order they blocked
 - ***Beware priority problems!***

Our Messaging Protocol

- `MboxRecv ()` will receive one message
 - Block if nothing buffered
- Messages must be received in order
- Blocked receivers must be kept in order
 - ***Beware priority problems here, too!***

Process X
(Priority 5)

Mailbox 1

Process X (Priority 5)

```
ID=MboxCreate()
```

```
# ID is 1
```

Mailbox 1

Process X (Priority 5)

```
ID=MboxCreate()  
  
# ID is 1  
  
fork1(pr=4) # A
```

Process A (Priority 4)

```
MboxRecv(1)  
# Block!
```


**Consumer
Queue**

Process
A




Mailbox 1

Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>




Mailbox 1


Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>




Mailbox 1


Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)	Process D (Priority 1)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	







Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)	Process D (Priority 1)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxSend(1, M1)</pre>







Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)	Process D (Priority 1)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxSend(1, M1) MboxSend(1, M2)</pre>







Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)	Process D (Priority 1)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D </pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxSend(1, M1) MboxSend(1, M2) MboxSend(1, M3)</pre>







Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)	Process D (Priority 1)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxSend(1, M1) MboxSend(1, M2) MboxSend(1, M3) quit()</pre>





Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)	Process D (Priority 1)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxSend(1, M1) MboxSend(1, M2) MboxSend(1, M3) quit()</pre>







Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>


***Who
should be
woken up
next?***

***Who
should
receive
M1?***






Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>






Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1) # Block!</pre>	<pre>MboxRecv(1) # Block!</pre>







Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D </pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1) # Block!</pre>




Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1) # Block!</pre>





Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1)</pre>



Mailbox 1

Process X (Priority 5)	Process A (Priority 4)	Process B (Priority 3)	Process C (Priority 2)
<pre>ID=MboxCreate() # ID is 1 fork1(pr=4) # A fork1(pr=3) # B fork1(pr=2) # C fork1(pr=1) # D</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1)</pre>	<pre>MboxRecv(1)</pre>

**Consumer
Queue**

Our Messaging Protocol

- We have a large pool of “message buffers”
 - Global pool, shared by all mailboxes
 - Allocate when you have a message
 - In `MboxSend()`, not `MboxCreate()`
 - Over-allocation is allowed
 - Buffering can fail if no free slots
 - Free in `MboxRecv()`

Our Messaging Protocol

- `MboxCondSend()` , `MboxCondRecv()` are **non-blocking** versions
 - Send or recv if possible
 - Return special value if you would have blocked

Messages

Mutexes & Signaling

Mutexes & Signaling

- A mailbox can be used (abused?) as a mutex
 - Initialize with size 1
 - `Send()` to lock, `Recv()` to unlock
 - Process A calls `Send()` while in CS
Process B calls `Send()` and blocks until A calls `Recv()`

ICA: Mutexes & Signaling

- A mailbox can be used (abused?) as a mutex
 - Initialize with size 1
 - `Send()` to lock, `Recv()` to unlock
 - Process A calls `Send()` while in CS
Process B calls `Send()` and blocks until A calls `Recv()`

Can it be reversed? Use `Recv()` to gain lock and `Send()` to release it? Discuss!

Mutexes & Signaling

- Soon, we'll be implementing syscalls
 - For example, call to `fopen` should block after the call until the the device responds
 - In-General: Block current process until something happens

ICA: Mutexes & Signaling

- Soon, we'll be implementing syscalls
 - For example, call to `fopen` should block after the call until the the device responds
 - In-General: Block current process until something happens

How to use mailboxes for sleep & wakeup?

Nicer interface than calling Phase 1 code

How to do it?

When to allocate mailboxes?

What can be shared, or not?