

---

# Semaphores

# Semaphores

---

- Concurrency Control Mechanisms, In General
- Semaphore Concepts
- Producer / Consumer

# Semaphores

---

## Concurrency Control Mechanisms, In General

# C.C. Mechanisms

---

- We have studied **locks** using **atomic instructions**
  - What are these two things?
- We have studied **messages / mailboxes**
- We noted that you could build a lock out of a messaging system
- **T,P,S**: Can you build a messaging system using locks?

# C.C. Mechanisms

---

- **General rule:**
- You can build any C.C. mechanism out of any other (might be ugly), if it has the two core requirements:
  - Mutex
  - Sleep / Wakeup

Two lesser requirements:

- Arbitrary # of processes
- No need for continual participation (“a process can go on vacation”)

# C.C. Mechanisms

---

## Mutex

- The system needs a way for one process to block another from having access to shared data

## Sleep / Wakeup

- System puts a process to sleep (no CPU) if it blocks for a long time

# C.C. Mechanisms

---

- **Messages with from Locks**
- Used a spinlock (and/or disabling interrupts) to control access to data structures (mutex)
- Changed process state, called dispatcher, when needed to block (sleep / wakeup)

# C.C. Mechanisms

---

- **Locks built from Messages**
- Read a message when you wanted a lock.  
Would block if no message. (mutex, sleep)
- Send a message when you release the lock  
(wakeup)



# Semaphores

---

## Semaphore Concepts

# Semaphores

---

- A **semaphore** is a concurrency control mechanism built around a *counter*
  - Can never go negative
  - Increment / decrement
- Designed by Edgar Dijkstra
- One of the first C.C. mechanisms  
(older than locks, maybe?)

# Semaphores

---

- **P ( )** (“Proberen:” Dutch for “try”) (also called **wait**)
  - If counter is zero, block until nonzero
  - Then decrement
- **V ( )** (“Verhogen:” Dutch for “increase”) (also called **post**)
  - Increment counter
  - If any **P ( )** s are blocked, unblock one

# Semaphores

---

- A **binary semaphore** is:
  - Initialized to 1
  - Never intended to increase beyond
- Can be used as a lock
- Optionally, can be simplified to a single bit

# Semaphores

---

- A **counting semaphore** is:
  - Initialized to any (non-neg) value
  - Can get large
- Used to count available resources
- Blocks if everything is busy

# ICA: Semaphores

---

Adapt our old concurrency example to use a semaphore. Use the semaphore to *protect the critical section*; do not use it to do the counting.

```
while True:
    load  x → $r1
    inc   $r1
    store $r1 → x
```

# Semaphores

---

```
s = init_semaphore(1)
```

```
while True:
```

```
    P(s)
```

```
    load    x → $r1
```

```
    inc     $r1
```

```
    store   $r1 → x
```

```
    V(s)
```

# Semaphores

---

```
s = init_semaphore(1)
```

```
while True:
```



```
P(s)
```

```
load    x → $r1
```

```
inc     $r1
```

```
store   $r1 → x
```



```
V(s)
```

Notice that we are treating the P() and V() operations as atomic ops – even though they are subroutines.

Probably, the functions use atomic instructions or locks inside, to provide this property.



# ICA: Write the Pseudocode

- Use a *counting semaphore* to simulate two types of concurrent processes: red processes need only one widget to do anything, but blue need two. (There can be many red and many blue processes.)
- The system starts with 5 widgets. Use the counting semaphore to keep track of how many are free.
- **Write** `start_red()`, `end_red()`, `start_blue()`, `end_blue()`

```
w = init_semaphore(5)
```

```
start_red() :
```

```
    P(w)
```

```
start_blue() :
```

```
    P(w)
```

```
    P(w)
```

```
end_red() :
```

```
    V(w)
```

```
end_blue() :
```

```
    V(w)
```

```
    V(w)
```

***What's the bug in this code?***

# Semaphores

---

- Semaphores are susceptible to deadlock, just like any other c.c. mechanism!

***What were the 4 conditions necessary for deadlock?***

# Semaphores

---

- Semaphores are susceptible to deadlock, just like any other c.c. mechanism!
- 4 conditions for deadlock:
  - Mutual exclusion
  - No pre-emption
  - Hold and wait
  - Circular wait

# Semaphores

---

- **Mutual Exclusion**
- Yes, a counting semaphore can allow multiple users at once
- But once a resource has been “allocated” with  $P()$ , no one else can use it
- Once all the resources have been claimed, all  $P()$  operations block

```
w = init_semaphore(5)
```

```
start_red() :    start_blue() :
```

```
    P(w)
```

```
    P(w)
```

```
    P(w)
```


```
end_red() :
```

```
    V(w)
```

```
end_blue() :
```

```
    V(w)
```

```
    V(w)
```



This is the only place in our code where Hold and Wait occurs.

Circular Wait happens when many processes block in  
`start_blue()`

# Producer / Consumer

---

- A classic form of process interaction is a producer / consumer
- One or more **producer** processes deliver blocks of data (messages, text, whatever) into a shared storage
  - Block when full
- One or more **consumer** processes consume the same blocks from the shared storage
  - Block when empty

# ICA: Producers and Consumers

---

- What is a practical use-case for the producers / consumers model?
- What problem(s) could this be used to solve with multiple processes?



# ICA: Producers and Consumers


---

- What is a practical use-case for the producers / consumers model?
- What problem(s) could this be used to solve with multiple processes?

PIZZA EXAMPLE!

Producer Threads

Consumer Threads



## Code as a Group:

- Use a pair of counting semaphores to model a producer / consumer system.
- Write `produce()`, `consume()`. Block the processes when necessary. (Assume that the system has only 10 data buffers, and it starts with all of them empty.)
- You may hand-wave the code needed to actually allocate the buffers and copy the data blocks; we are focused on concurrency control.

```
f = init_semaphore(0)
e = init_semaphore(10)
```

```
produce() :
    P(e)
    ...data...
    V(f)
```

```
consume() :
    P(f)
    ...data...
    V(e)
```