

# Memory Management

CS 452

Part 1

# Abstracting a CPU

- Have discussed at length the concept of a process
- Process = an abstraction of a CPU
  - Allows processes to not have to worry about the details of sharing the CPU hardware, can act as if they “own” it
- How do we handle *memory*?

# A Simple Scenario

- Computer has exactly 16 bytes of memory
- Addresses 0x0 through 0xf
- Let's consider how this would work for a few simple scenarios

# Scenario: One Program

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
load temp mem(0x5)
print(temp)
```

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

*\*\* Note: This is an over-simplification - we would need memory for the instructions, vars, etc too*

# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
load temp mem(0x5)
print(temp)
```

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
load temp mem(0x4)
print(temp)
```

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

*\*\* Note: This is an over-simplification - we would need memory for the instructions, vars, etc too*

# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
load temp mem(0x5)
print(temp)
```

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
load temp mem(0x4)
print(temp)
```

*What happens if a context switch from A to B happens half-way through execution?*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
load temp mem(0x5)
print(temp)
```

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
load temp mem(0x4)
print(temp)
```

*What are possible solutions?*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
load temp mem(0x5)
print(temp)
```

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
load temp mem(0x4)
print(temp)
```

***Solution 1: Each program gets full memory, save to disk when CSed away from, load when CSed back to***

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	



# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
load temp mem(0x5)
print(temp)
```

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
load temp mem(0x4)
print(temp)
```

***Solution 2: Programs only get as much memory as they need. Have a limited range. Problems??***

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	


# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```



## Problems:

*(A) How to know address range in advance?*

*(B) What if there isn't enough space?*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# Scenario: Two Programs

## Program A

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**BASE = 0x0**

## Program B

```
i = 0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x8**

## *Solution for A:*

*Each is written to assume it starts at 0x0, is given an offset address*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# Dynamic Relocation

- This technique is known as **dynamic relocation**
- Either when program loaded or as instructions execute, dynamically change the addresses if need be
- **BASE** Register: The offset for all addresses the program uses
- **LIMIT** Register: Represents the max address
  - Combined gives us an **Address Space!**
  - Each program can use any address it wants between 0x0 and **LIMIT**. Kernel / CPU handles the offset.

# Scenario: Three Programs

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**BASE = 0x0**

## Program B

```
i = 0x0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x5**

## Program C

```
c = 25
store c mem(0x0)
store c mem(0x1)
store c mem(0x2)
store c mem(0x3)
```

**BASE = 0xd**

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# Scenario: Three Programs

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**BASE = 0x0**

## Program B

```
i = 0x0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x5**

## Program C

```
c = 25
store c mem(0x0)
store c mem(0x1)
store c mem(0x2)
store c mem(0x3)
```

**BASE = 0xd**

*What is the problem?*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# Scenario: Three Programs

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**BASE = 0x0**

## Program B

```
i = 0x0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x5**

## Program C

```
c = 25
store c mem(0x0)
store c mem(0x1)
store c mem(0x2)
store c mem(0x3)
```

**BASE = 0xd**

*Not enough space!*

*We are one byte of memory short.*

*Program C either needs to wait (could starve)  
or could copy to memory*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

# A Different Scenario

- Don't change anything, except:
  - Computer has more memory 64kb bytes of memory
  - Addresses 0x0 through 0xffff
- Now we have PLENTY of memory to run tons of programs! (😄)
- But.... there still is a problem



# Scenario: Plenty of Memory

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

.... plenty more ....

# Scenario: Plenty of Memory

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**BASE = 0x0**

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

.... plenty more ....

# Scenario: Plenty of Memory

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

## Program B

```
i = 0x0
while i <= 0x8:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x0**

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

.... plenty more ....

# Scenario: Plenty of Memory

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**BASE = 0x0**

## Program B

```
i = 0x0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x5**

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

.... plenty more ....

# Scenario: Plenty of Memory

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**malloc(input())**

**BASE = 0x0**

## Program B

```
i = 0x0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x5**

*Process A needs more memory, what to do?*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

.... plenty more ....

# Scenario: Plenty of Memory

## Program A

```
i = 0
while i <= 0x4:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

**malloc(input())**

**BASE = 0x0**

## Program B

```
i = 0x0
while i <= 0x7:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

**BASE = 0x5**

*Process A needs more memory, what to do?*

*- Move it to a new location?*

*- Shift process B down?*

Address	Value
0x0	
0x1	
0x2	
0x3	
0x4	
0x5	
0x6	
0x7	
0x8	
0x9	
0xa	
0xb	
0xc	
0xd	
0xe	
0xf	

.... plenty more ....

# Not Satisfying

- Each of the proposed solutions has one or more weaknesses
- Need a mechanism that gives each process its own **abstract memory space**, while still being relatively efficient.
- An Answer....

# Virtual Memory with Page Tables

- **Virtual Memory** is the technique used in many modern OSs
- Each process gets to act as if it owns a large memory space completely to itself.
  - EG, on a 64-bit system, addresses `0x000000000000` to `0xffffffffffff`
- A combination of the OS kernel and the hardware manage the mapping between each process's **virtual address space** and the **physical address space** of the RAM.


*NOTE: No need for BASE / LIMIT anymore!*

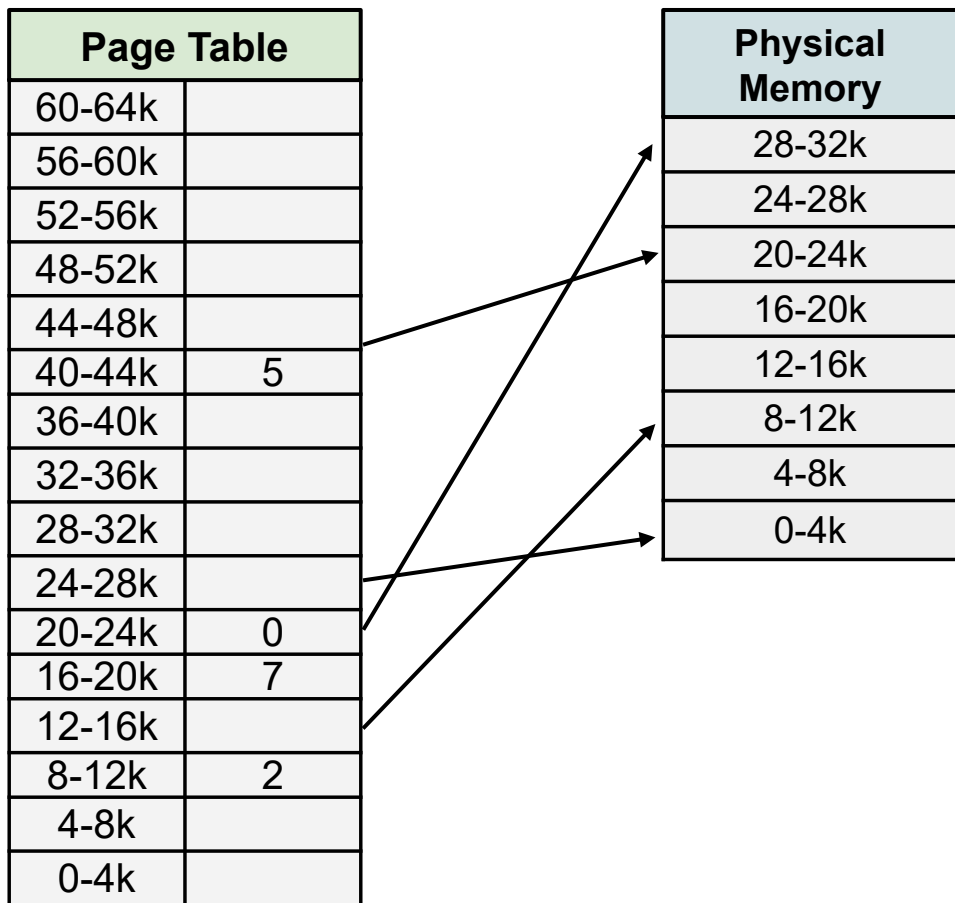


# Virtual Memory with Page Tables

- The **page table** is a key structure to allow for this to happen with reasonable efficiency
- The OS keeps a **page table** for each process
- Maps virtual addresses to physical ones
- Divides virtual memory spaces into pages (for example, 4k chunks of bytes)

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	1
20-24k	X
16-20k	X
12-16k	5
8-12k	X
4-8k	X
0-4k	0






# New Scenario

- Single-core computer
- Has 32k of physical memory
- Each process gets a **virtual address space** of 64k
- Page size is 4k
  - 8 physical page slots, 16 virtual pages per process

# Program A

```
i = 0
while i <= 54k:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X




## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

# Program A

```
i = 0
while i <= 54k:
    load r1 mem(i)
    r1++
    store r1 mem(i)
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	0



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



# Program A

```
i = 0
```

```
while i <= 54k:
```


```
    load r1 mem(i)
```

```
    r1++
```

```
    store r1 mem(i)
```

```
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	1
0-4k	0



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



# Program A

```
i = 0
```


```
while i <= 54k:
```

```
    load r1 mem(i)
```

```
    r1++
```

```
    store r1 mem(i)
```

```
    i++
```



Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	2
4-8k	1
0-4k	0

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k





# Program A

```
i = 0
```


```
while i <= 54k:
```

```
    load r1 mem(i)
```

```
    r1++
```

```
    store r1 mem(i)
```

```
    i++
```



Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	2
4-8k	1
0-4k	0

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

...

# Program A

```
i = 0
```

```
while i <= 54k:
```

```
    load r1 mem(i)
```

```
    r1++
```


```
    store r1 mem(i)
```

```
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	2
4-8k	1
0-4k	0

*What happens next?*

...



Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

# Storing Pages on Disk

- Sometimes, need to store pages on disk (hard drive, SSD, etc)
  - Virtual memory space often is larger than physical memory
  - Even if that was not the case, multiple processes
- These are stored / loaded into a **swap partition** or **swap file**
  - When page evicted, save if need be
  - When page table has miss, load from swap file
- See `$ cat /proc/PID/maps`

# Program A

```
i = 0
```

```
while i <= 54k:
```


```
    load r1 mem(i)
```

```
    r1++
```

```
    store r1 mem(i)
```

```
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	2
4-8k	1
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

**Contents of Virtual Page 0 saved to disk for (possible) future use**

**← Notice: This is now invalid!**

# Program A

```
i = 0
```

```
while i <= 54k:
```

```
    load r1 mem(i)
```


```
    r1++
```

```
    store r1 mem(i)
```

```
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	2
4-8k	X
0-4k	X

Contents of Virtual Page  
1 saved to disk for  
(possible) future use



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

Notice: This is now invalid!

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Context  
Switch!

## Program B

`i = 20k`

`while i <= 40k:`

`load r1 mem(i)`


`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	2
4-8k	X
0-4k	X

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Context  
Switch!

## Program B

`i = 20k`


`while i <= 40k:`

`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`



Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	2
4-8k	X
0-4k	X

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

*How to  
choose next  
physical  
memory  
location to  
use?*

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	3
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k





## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	4
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

### Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	5
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



Old Content from  
Proc A goes to disk

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	6
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	5
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	1
32-36k	0
28-32k	7
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`



`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



Old Content from  
Proc A goes to disk

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	7
36-40k	1
32-36k	0
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	0
40-44k	7
36-40k	1
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	2
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



**Old Content from  
Proc A goes to disk**

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



## Program A

`i = 0`

`while i <= 54k:`

`load r1 mem(i)`

`r1++`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



Old Content from  
Proc B goes to disk

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

`i = 0`

`while i <= 54k:`

`. . . .`

`load r1 0x10`

## Program B

`i = 20k`


`while i <= 40k:`

`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`



Disk

Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

*What happens if this is the next instruction executed?*

## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

# Non-Sequential

- Entries in the page table do not have to be sequential
- Programs can have “sequential” access patterns, but also “random” access patterns as well

## Program A

`i = 0`

`while i <= 54k:`

`. . . .`

`load r1 0x10`

**Invalid**  
**!**

Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

`i = 20k`

`while i <= 40k:`


`load r1 mem(i)`

`r1 = r1 * 2`

`store r1 mem(i)`

`i++`

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

**Save!**

## Program A


```
i = 0
while i <= 54k:
    . . . .
    load r1 0x10
```

Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

## Program B

```
i = 20k
while i <= 40k:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```


Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	3
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

Load  
Page



## Program A


```
i = 0
while i <= 54k:
    . . . .
    load r1 0x10
```

Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	3

## Program B


```
i = 20k
while i <= 40k:
    load r1 mem(i)
    r1 = r1 * 2
    store r1 mem(i)
    i++
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



## Program A

`i = 0`

`while i <= 54k:`

`. . . .`

`load r1 0x10`


## Program B

`i = 20k`

`while i <= 40k:`


`. . . .`

`store r2 0x5208`



Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	3

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

## Program A

```
i = 0
while i <= 54k:
    . . . .
    load r1 0x10
```


Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	3

## Program B

```
i = 20k
while i <= 40k:
    . . . .
    store r2 0x5208 // Dec 21000
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

Invalid  
!



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

Save  
and  
Load



## Program A


```
i = 0
while i <= 54k:
    . . . .
load r1 0x10
```

Page Table	
60-64k	X
56-60k	X
52-56k	2
48-52k	1
44-48k	0
40-44k	7
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	3

## Program B

```
i = 20k
while i <= 40k:
    . . . .
store r2 0x5208 // Dec 21000
```

Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	6
32-36k	5
28-32k	X
24-28k	X
20-24k	4
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



## Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



# Page Size

*Is a 4k Page size a good idea?*

*What happens if we make the pages size larger? 8k? 16k?*

*What about smaller?*

# Page Table Lookup

- How does a page table lookup work?
- The incoming address is split up in to two components a **Page Table Index** (PTI) and an **Offset**.
- PTI used to determine slot in page table
- Offset is appended to the value stored in page table to get physical address

**Present / Absent bit**

**Page Table index  
(Decimal)**

**Physical Page  
Number**

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

## Incoming Virtual Address

0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

	Page Table	
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

## Incoming Virtual Address

0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

## Incoming Virtual Address

0	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---





1	1	0													
---	---	---	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

### Page Table

15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1



[illegible]

1	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

	Page Table	
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1



## Incoming Virtual Address

1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

*What about this scenario?  
Trace the steps*

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

Outgoing Physical Address

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

## Incoming Virtual Address

1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

## Incoming Virtual Address

1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



**Absent!**  
**Must be swapped in**

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	000	0
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	010	1

## Incoming Virtual Address

1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



**Absent!**  
**Must be swapped in**

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	010	1
9	<del>101</del>	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	000	0

## Incoming Virtual Address

1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	010	1
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	000	0

0	1	0													
---	---	---	--	--	--	--	--	--	--	--	--	--	--	--	--

## Outgoing Physical Address

## Incoming Virtual Address

1	0	1	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



0	1	0	0	0	0	0	0	0	0	0	0	0	1	0	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

## Outgoing Physical Address

Page Table		
15	000	0
14	000	0
13	000	0
12	000	0
11	111	1
10	010	1
9	101	1
8	000	0
7	000	0
6	000	0
5	011	1
4	100	1
3	000	1
2	110	1
1	001	1
0	000	0

# Page Table Entry

- Each entry needs some information, in addition to present bit and the frame number
  - Protection: Permission to Read? Write? Exec?
  - Modification: Has this data been modified?
  - Referenced: Has this page been referenced?
  - Can this be cached (for optimization reasons)?

<i>extra</i>	caching en / dis	refd	mod	protect	present / absent	Frame Number
--------------	---------------------	------	-----	---------	------------------------	--------------

# Page Tables - Two Implementations

- Have an in-hardware page table in the MMU (on or near the CPU). Whenever a context switch happens, load the processes page table into hardware, have MMU do all instruction translations
- Have a register in the CPU that points to the current processes page table. Update this register whenever there is a context switch happens, use MMU+Memory to translate addresses



# Page Tables - Two Implementations

- Have an in-hardware page table in the MMU (on or near the CPU). Whenever a context switch happens, load the processes page table into hardware, have MMU do all instruction translations
- Have a register in the CPU that points to the current processes page table. Update this register whenever there is a context switch happens, use MMU+Memory to translate addresses

***What are the pros and cons of each?***

# Page Tables - Two Implementations

- Dedicated PT Hardware
  - Pro: Once loaded, fast!
  - Cons: Context switching expensive, could be huge table
- Memory PT
  - Pro: Super cheap to context switch
  - Con: Slower overall execution

***The cons in both cases are NOT acceptable for performance!***

# Starting a Process

- When a process first begins, how does the OS know where to get the memory from?
  - instructions, space for global vars, etc
- From the ELF file!
- Includes VMA (Virtual Memory Address) Information



# Fork and Exec

- What to do when we want to create a new process?
- In a UNIX System, call **fork()**
  - Create new entry in process table, get new PID, etc
  - What do do with all the memory?

# Fork and Exec

- What to do when we want to create a new process?
- In a UNIX System, call **fork()**
  - Create new entry in process table, get new PID, etc
  - What do do with all the memory?
- Gets an exact memory copy!
- Does not need to copy it all at once, use Copy-On-Write instead.

# Fork and Exec

- Often, After a fork occurs, there is a follow-up call to **exec**
- Exec replaces the memory for the current process
- Used in tandem to create new processes with new memories

# Page Tables - Two Implementations

*What hardware could be added to a memory-based page table system to provide a speed boost?*



# Use Dedicated Hardware + Memory

- Key observation: Most programs make a large number of references to a small number of pages, not the other way around
- Storing the entire page table in hardware is likely super excessive
- Instead, store a subset of page table entries in hardware
- **Translation Lookaside Buffers!**

# Translation Lookaside Buffers

A **Translation Lookaside Buffer** (TLB) is a relatively small array-like hardware structure that stores a small number of page table entries. Ideally, ones that are actively / often being used by the CPU.

For example, TLB could store 32 PT entries. If former observation is true, the majority of memory references will use this instead of needing to look at in-memory PT.



# Program A

```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main()
```



Page Table	
60-64k	X
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

TLB				
Valid	Virt Page	Mod	Prot	Page Frame

Physical Memory
28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

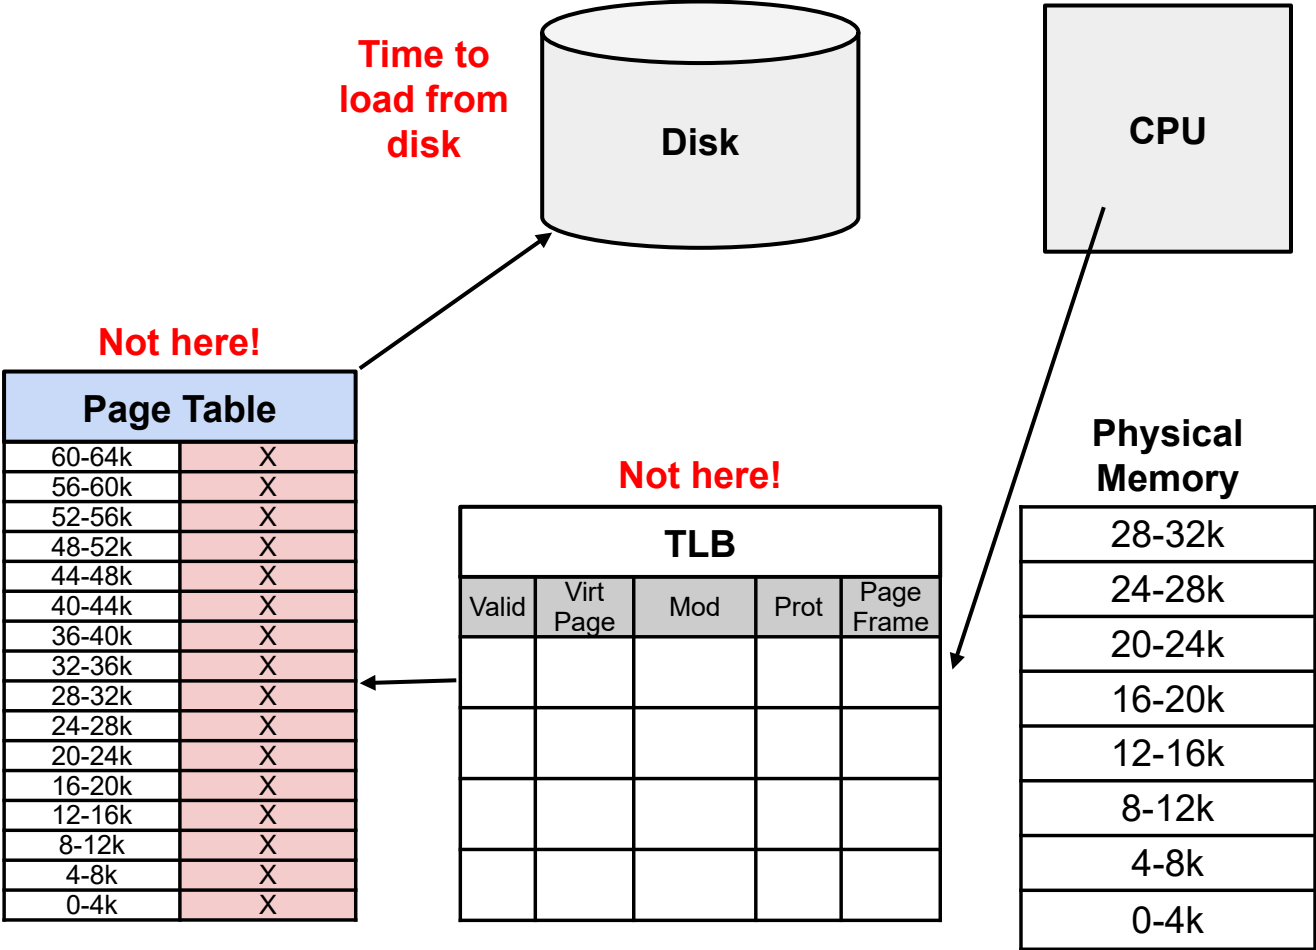
# Program A

```
print('hi')
let arr;

def calc(x):
  i = 0;
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main()
```

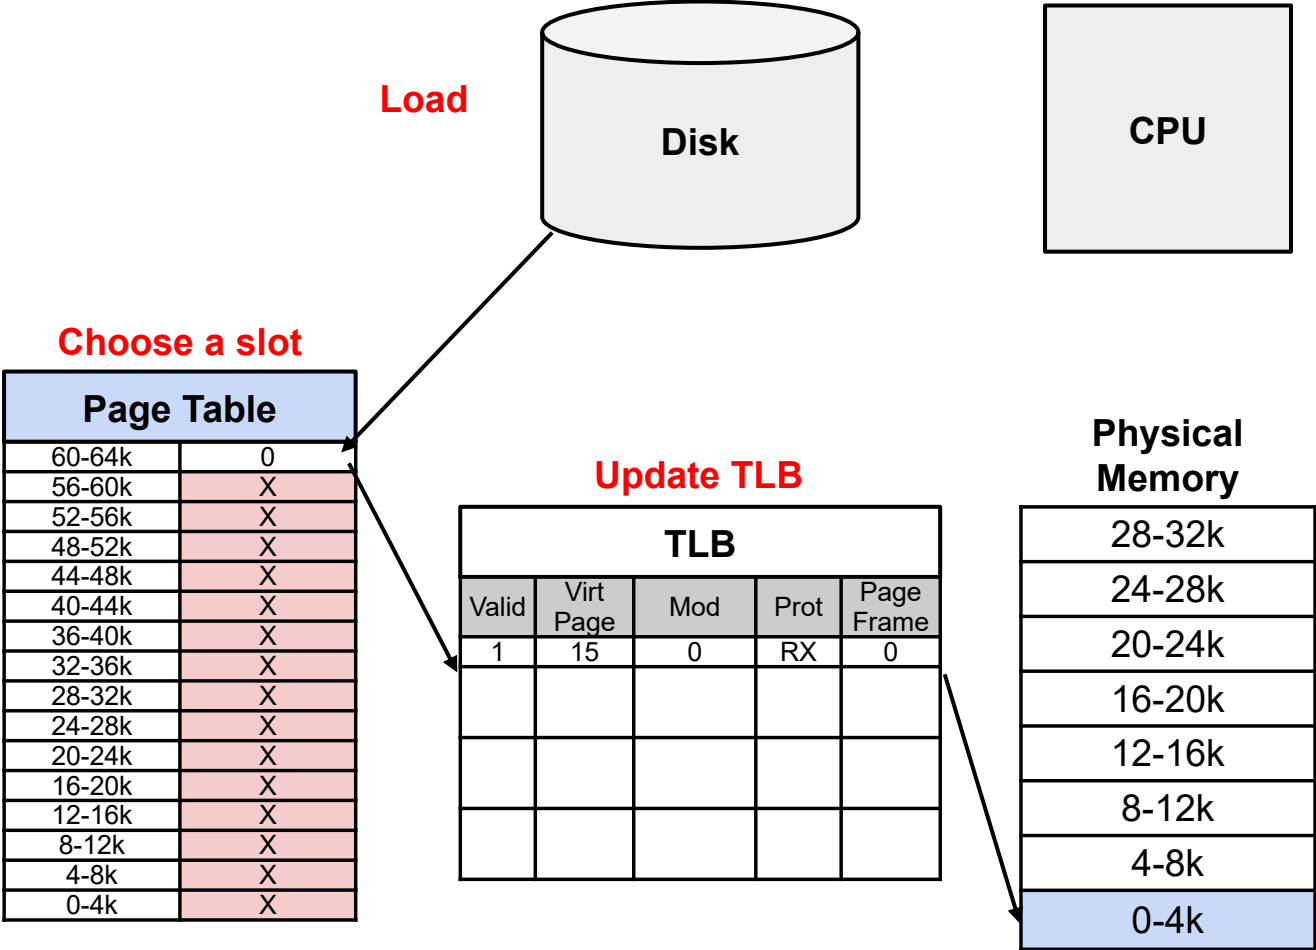


# Program A

```
print('hi')  
let arr;
```

```
def calc(x):  
    i = 0;  
    while i < 10000:  
        arr[i%100] = x  
        i++
```

```
def main():  
    arr = alloc(100)  
    calc(5)  
    calc(10)  
  
main()
```



# Program A

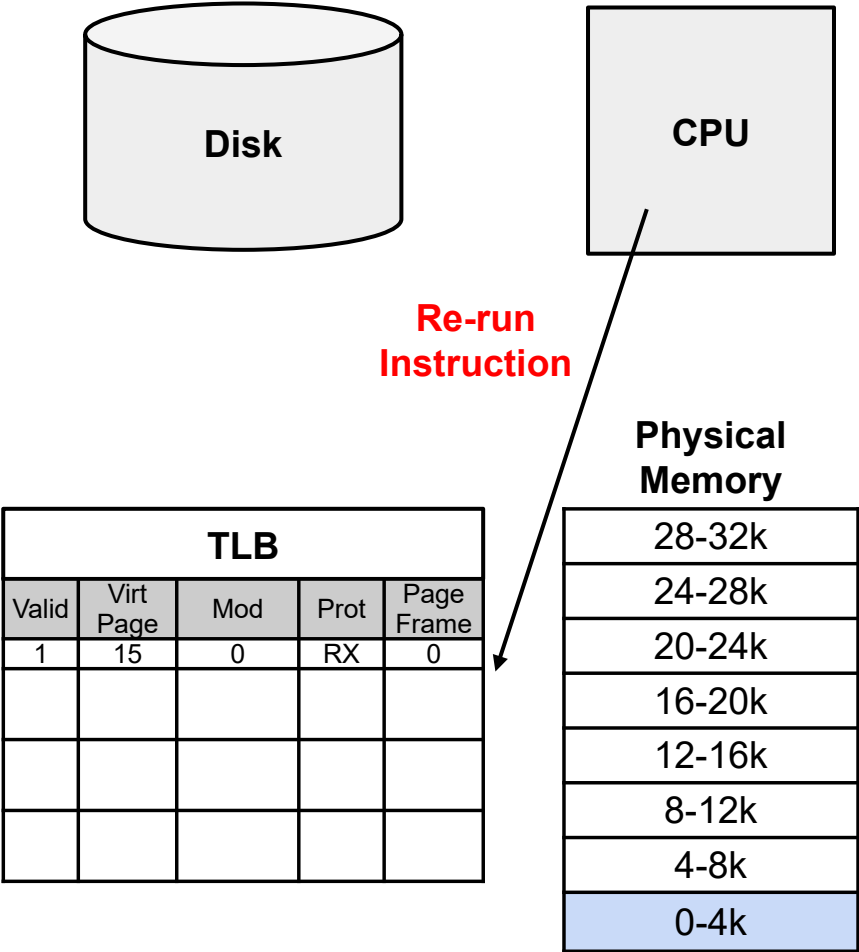
```
print('hi')  
let arr;
```

```
def calc(x):  
    i = 0;  
    while i < 10000:  
        arr[i%100] = x  
        i++
```

```
def main():  
    arr = alloc(100)  
    calc(5)  
    calc(10)
```

```
main()
```

Page Table	
60-64k	0
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



# Program A

```
print('hi')
let arr;



def calc(x):
    i = 0;
    while i < 10000:
        arr[i%100] = x
        i++

def main():
    arr = alloc(100)
    calc(5)
    calc(10)

main()
```

This probably all fit in that one page. Now, all instructions loaded into memory.

Page Table	
60-64k	0
56-60k	X
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0

Physical Memory	
28-32k	
24-28k	
20-24k	
16-20k	
12-16k	
8-12k	
4-8k	
0-4k	

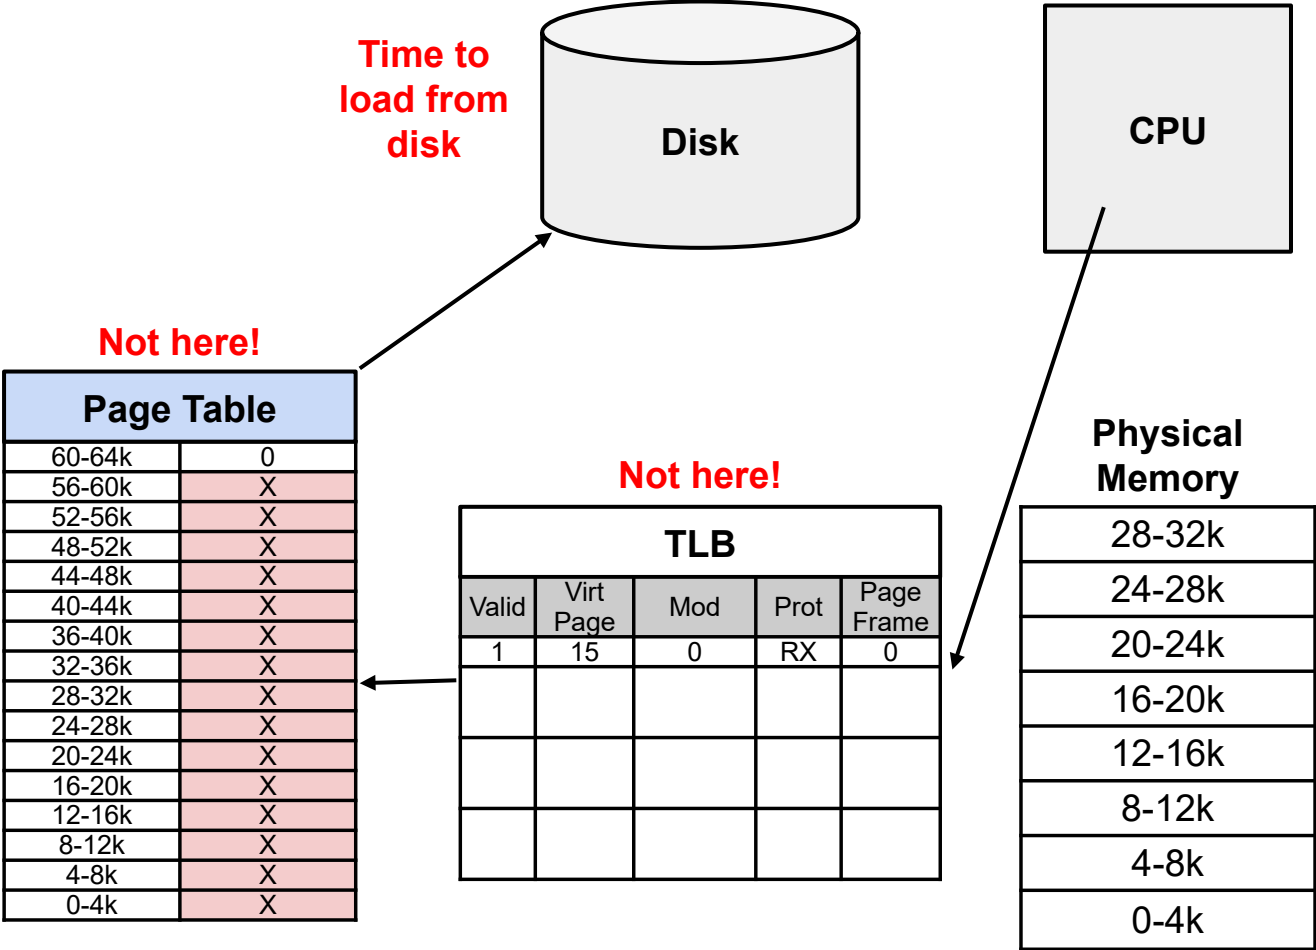
# Program A

```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main() ←
```





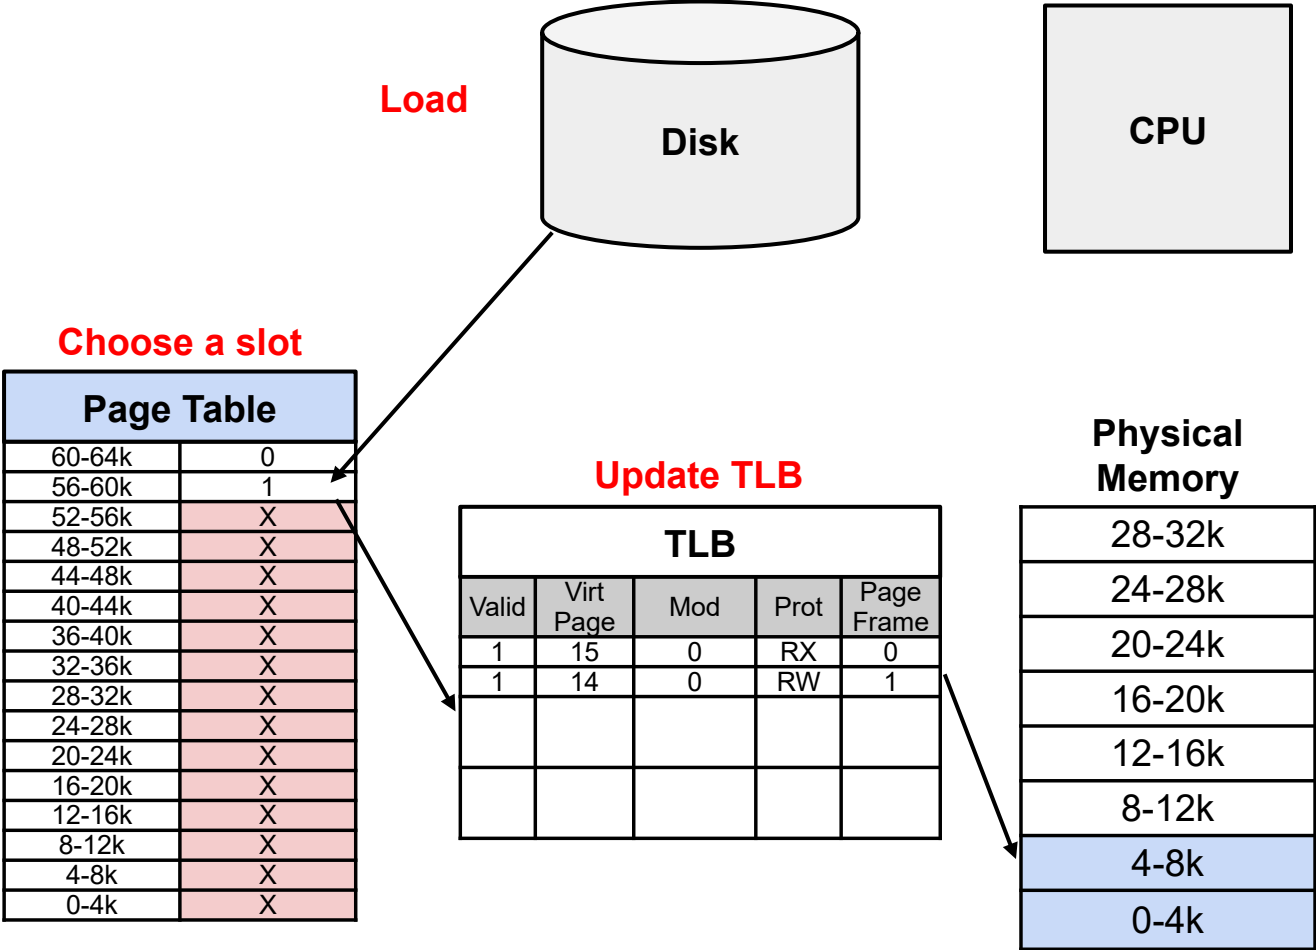
# Program A

```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main() ←
```



# Program A



```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main() ←
```


Is should be enough  
space for the stack as  
this program executes!



Page Table	
60-64k	0
56-60k	1
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	X
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0
1	14	0	RW	1

Physical Memory	
28-32k	
24-28k	
20-24k	
16-20k	
12-16k	
8-12k	
4-8k	
0-4k	



# Program A



```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main()
```

Page Table	
60-64k	0
56-60k	1
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	2
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X



Update TLB

TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0
1	14	0	RW	1

Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k



# Program A

```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main()
```



Choose a slot

Page Table	
60-64k	0
56-60k	1
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

Update TLB

TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0
1	14	0	RW	1
1	7	0	RW	4

Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

# Program A



```
print('hi')
let arr;

def calc(x):
  i = 0:
  while i < 10000:
    arr[i%100] = x
    i++

def main():
  arr = alloc(100)
  calc(5)
  calc(10)

main()
```

This should fit entire array!



Choose a slot

Page Table	
60-64k	0
56-60k	1
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

Update TLB

TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0
1	14	0	RW	1
1	7	0	RW	4

Physical Memory

28-32k
24-28k
20-24k
16-20k
12-16k
8-12k
4-8k
0-4k

# Program A

```
print('hi')
let arr;

def calc(x):
    i = 0:
    while i < 10000:
        arr[i%100] = x
        i++


def main():
    arr = alloc(100)
    calc(5)
    calc(10)

main()
```

From here on out, should be mostly TLB hits as CPU executes instructions.

Page Table	
60-64k	0
56-60k	1
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0
1	14	0	RW	1
1	7	0	RW	4



Physical Memory	
28-32k	
24-28k	
20-24k	
16-20k	
12-16k	
8-12k	
4-8k	
0-4k	

# Program A

```
print('hi')
let arr;


def calc(x):
    i = 0;
    while i < 10000:
        arr[i%100] = x
        i++

def main():
    arr = alloc(100)
    calc(5)
    calc(10)

main()
```

## What about Context Switches?

## What happens to the TLB?



Page Table	
60-64k	0
56-60k	1
52-56k	X
48-52k	X
44-48k	X
40-44k	X
36-40k	X
32-36k	X
28-32k	4
24-28k	X
20-24k	X
16-20k	X
12-16k	X
8-12k	X
4-8k	X
0-4k	X

TLB				
Valid	Virt Page	Mod	Prot	Page Frame
1	15	0	RX	0
1	14	0	RW	1
1	7	0	RW	4

Physical Memory	
28-32k	
24-28k	
20-24k	
16-20k	
12-16k	
8-12k	
4-8k	
0-4k	

# Hardware vs Software Handling

- Could have the MMU hardware handle all TLB misses directly
- Or, could have an interrupt occur whenever TLB miss happens, have kernel handle the issue



# Hardware vs Software Handling

- Could have the MMU hardware handle all TLB misses directly
- Or, could have an interrupt occur whenever TLB miss happens, have kernel handle the issue

***What are the pros and cons of each?***

# TLB Hits and Misses

- **Hit:** A virtual memory address is in TLB
- **Soft Miss:** A virtual memory address is not in TLB, but is in memory
- **Hard Miss:** A virtual memory address lookup is not in TLB or in memory, go to disk instead

# TLB Hits and Misses

- **Hit:** A virtual memory address is in TLB
- **Soft Miss:** A virtual memory address is not in TLB, but is in memory
- **Hard Miss:** A virtual memory address lookup is not in TLB or in memory, go to disk instead

*Performance of each?*

# Page Replacement Algorithms

When it comes to to evict a page out of physical memory and replace with another, how to make this choice?

What algorithms can you think of that might work well?

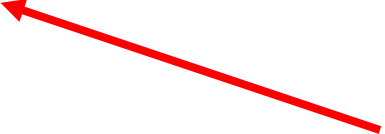
# MMAP

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/mman.h>
int main() {
    int fd = open("/bin/bash", O_RDONLY);
    printf("fd %d\n", fd);
    char *buf = mmap(NULL, 4096,
                     PROT_READ,
                     MAP_PRIVATE,
                     fd, 0);
    printf("The first 8 bytes are:\n");
    for (int i=0; i<8; i++)
        printf(" %d: 0x%02x\n", i, buf[i]);
    return 0;
}
```

# MMAP

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/mman.h>
int main() {
    int fd = open("/bin/bash", O_RDONLY);
    printf("fd %d\n", fd);
    char *buf = mmap(NULL, 4096,
                     PROT_READ,
                     MAP_PRIVATE,
                     fd, 0);

    printf("The first 8 bytes are:\n");
    for (int i=0; i<8; i++)
        printf(" %d: 0x%02x\n", i, buf[i]);
    return 0;
}
```



open returns a  
file descriptor,  
needed for using  
mmap

# MMAP

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/mman.h>
int main() {
    int fd = open("/bin/bash", O_RDONLY);
    printf("fd %d\n", fd);
    char *buf = mmap(NULL, 4096,
                     PROT_READ,
                     MAP_PRIVATE,
                     fd, 0);

    printf("The first 8 bytes are:\n");
    for (int i=0; i<8; i++)
        printf(" %d: 0x%02x\n", i, buf[i]);
    return 0;
}
```

**NULL for  
address, one  
page (4094)**



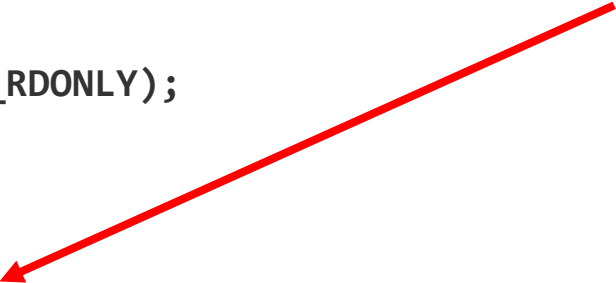
**Create a read-  
only page**



# MMAP

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/mman.h>
int main() {
    int fd = open("/bin/bash", O_RDONLY);
    printf("fd %d\n", fd);
    char *buf = mmap(NULL, 4096,
                     PROT_READ,
                     MAP_PRIVATE,
                     fd, 0);
    printf("The first 8 bytes are:\n");
    for (int i=0; i<8; i++)
        printf(" %d: 0x%02x\n", i, buf[i]);
    return 0;
}
```

This copy is only for this process to use, not shared



Use the fd file, start at byte 0



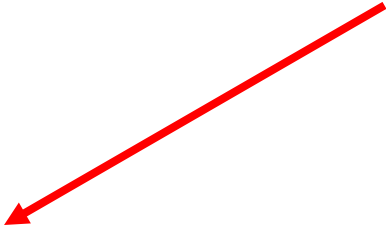


# MMAP

```
#include <stdio.h>
#include <fcntl.h>
#include <sys/mman.h>
int main() {
    int fd = open("/bin/bash", O_RDONLY);
    printf("fd %d\n", fd);
    char *buf = mmap(NULL, 4096,
                     PROT_READ | PROT_WRITE,
                     MAP_SHARED,
                     fd, 0);

    printf("The first 8 bytes are:\n");
    for (int i=0; i<8; i++)
        printf(" %d: 0x%02x\n", i, buf[i]);
    return 0;
}
```

**Can also  
configure to be  
shareable!**



**Other processes  
forked from this  
one will be able  
to see the same  
page**

# MMap and MProtect Example

1. Copy `/tmp/m_sharing_protecting.c` to a directory you own
2. Read the code. Can you tell me what it does?
3. Create a file named `testing.txt` in the same directory, compile, run
4. What will be in `testing.txt` after it runs?

# MMap and MProtect Example

What happens is the parent processes tries to write something to the shared buffer?

# Virtual Memory

1. Copy `/tmp/m_small.c` to a directory you own
2. Read the code. Can you tell me what it does?

# Virtual Memory

1. Copy `/tmp/m_small.c` to a directory you own
2. Read the code. Can you tell me what it does?
3. Now, do the same for `/tmp/m_large.c`. What is the difference in the code? Performance difference?

# Virtual Memory

1. Copy `/tmp/m_small.c` to a directory you own
2. Read the code. Can you tell me what it does?
3. Now, do the same for `/tmp/m_large.c`. What is the difference in the code? Performance difference?
4. Now, do the same for `/tmp/m_multiple_passes.c`.

# Checking for Page Table Misses

- How to look at the page table misses?
  - min  $\approx$  soft miss and maj  $\approx$  hard miss

```
$ ps -eo min_flt,maj_flt,cmd | grep -e m_reuse -e MINFL
```

# Checking for Page Table Misses

- How to look at the page table misses?
  - `min`  $\approx$  soft miss and `maj`  $\approx$  hard miss

```
$ ps -eo min_flt,maj_flt,cmd | grep -e m_reuse -e MINFL
```

- The `min_flt` and `maj_flt` can be viewed for whatever running processes you want.



# Attribution

Several examples in these slides were inspired by ones from  
“Modern Operating Systems” by Andrew Tanenbaum